# Applied Mathematics for Deep Learning I - Term Project

Group Members: Andre Graham (101555540), Adnan Sheik (101549676), Majed Qarmout (101565483), Yonotan Girma (101532284), Younes Bensassi-nour (101542854)

https://github.com/mqarmout/amazon\_sentiment\_analysis/tree/main

## 1. The Business Problem: Understanding Customer Reviews

Understanding customer sentiment is essential for businesses to effectively evaluate the public perception surrounding their products. Though this can be easily gauged on a platform such as Amazon that includes the user's rating out of five in the customer's review, platforms such as YouTube and Twitter provide more of a challenge since there is no inclusion of a rating. For companies to gain a comprehensive understanding of public perception surrounding their products, it is vital that they analyze customer sentiment across all platforms. Our model aims to determine the sentiment of unstructured text, facilitating companies to better understand public perception of their products. This will allow them to tailor their marketing strategies, improve their current and future product offerings, and enhance customer satisfaction.

#### 2. About the Data

The dataset used to train our model is the "Amazon Product Review" dataset from Kaggle. It has 10 features: Id, ProductId, UserId, ProfileName, HelpfulnessNumerator, HelpfulnessDenominator, Score, Time, Summary, and Text. For the purposes of our model, the two key variables are Score, which will be used to create the sentiment labels, and Text, which will be used as the input for our model. This is because, for the use case of our model, the other features will not be available as inputs. The Text column contains the text left in the review. The data was gathered through web scraping; however, the exact web scraping method is unknown, and therefore there may be a bias toward more popular products.

### 3. Our Approach

Our team's approach is to fine-tune an existing language model with product reviews. The goal is to leverage an existing sentiment analysis model that

already contains a deep understanding of natural language. We chose to use the RoBERTa model (robustly optimized BERT pretraining approach), which we accessed through the Hugging Face Transformers library. The issue with using a pretrained model such as RoBERTa is that it is trained on a variety of data, many of which are structured significantly differently from a product review on social media. We account for this by fine-tuning the model using the Amazon product reviews data, with the goal of increasing the model's accuracy when evaluating product reviews. We start by fine-tuning the "distilroberta-base" model by freezing all its layers and adding a feedforward network to the end of the model. Throughout this process, we fine-tune the hyperparameters, in total creating 12 potential models. We then compare this to fine-tuning the last 3 encoding layers of the "distilroberta-base" model, which has 6 encoding layers. Next, we pick the model with the best accuracy and create an inference pipeline. Finally, we list the limitations of our model and have reflections for each team member.

## 4. Show and plot dataset and apply data analysis on dataset

One of the most important aspects of the dataset to understand is the distribution of the score. The score is the rating out of 5 that the user gave the product and will be used to create the labels to train our model. Figure 1 shows the distribution of the scores.

```
In [2]: import pandas as pd
        import numpy as np
        import matplotlib.pyplot as plt
        # Loading the dataset
        df = pd.read csv('Data/Reviews.csv')
        # Setting plot size
        plt.figure(figsize=(12, 6))
        # Plotting the histogram
        plt.hist(df['Score'], bins=5, edgecolor='black')
        # Set up the labels and axis
        plt.title('Figure 1: Score Distribution')
        plt.xlabel('Score')
        plt.ylabel('Frequency')
        plt.xticks(range(1, 6))
        # Show the plot
        plt.show()
```

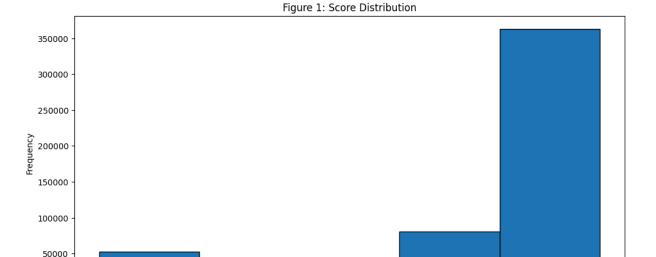


Figure 1 shows that the distribution of the scores is not uniform, which will have to be taken into account during preprocessing. This means that when we are creating our labels, we must either remove samples from categories with higher frequencies or add samples to the categories with low frequencies. Given we have limited computational resources, the more reasonable approach is to reduce the size of the dataset rather than increase it.

3 Score

Next, we visualize a word cloud of the text in the reviews. This will give us a general understanding of the text content of the reviews, as well as reveal any problematic patterns found in the text. This can be seen in figure 2.

```
In [3]: from wordcloud import WordCloud

# Word cloud for the text of the user reviews
text = ' '.join(df['Text'].astype(str))
wordcloud = WordCloud(width=800, height=400, background_color='white').gener

# Set the figure size for better readability
plt.figure(figsize=(12, 6))

# Plot the wordcloud
plt.imshow(wordcloud, interpolation='bilinear')

# Set up axis and labels
plt.axis('off')
plt.title('Figure 2: Word Cloud of Reviews', fontsize = 18)

# Show the plot
plt.show()
```

Figure 2: Word Cloud of Reviews



Figure 2 immediatly reveals a problem with the text in our dataset, which is the existence of tags. This is the reason why "br" is written twice in such large letters. It also makes sense that words such as love, favourite, and treat are all included in this word cloud given the skewed distribution of the ratings. These two issues will be addressed in our data preprocessing.

Finally, Figure 3 shows the distribution of the lengths of the reviews in the dataset.

```
In [4]: # Calculate the number of words in each review
    df['ReviewLength'] = df['Text'].apply(lambda x: len(x.split()))

# Set the figure size
    plt.figure(figsize=(12, 6))

# Plot the histogram
    plt.hist(df['ReviewLength'], bins=50, color='skyblue', edgecolor='black')

# Add titles and labels
    plt.title('Figure 3: Distribution of Review Lengths', fontsize=18)
    plt.xlabel('Number of Words', fontsize=14)
    plt.ylabel('Frequency', fontsize=14)

# Add grid lines for easier interpretation
    plt.grid(axis='y', alpha=0.75)

# Show the plot
    plt.show()
```

Figure 3: Distribution of Review Lengths

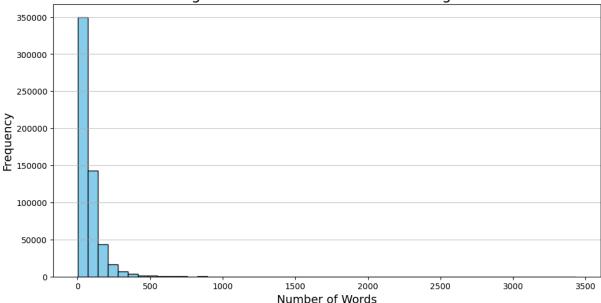


Figure 3 shows that that the distribution of the lengths of reviews is heavily right skewed. Since the goal of the model is to gain insight into the sentiment of customers posting on social media posts which tend to be quite short (think YouTube comments or Tweets), it is reasonable to conclude that the lengths of the text is representitive of the content that the model will be used on.

Though there are additional features included in the reviews dataset, the goal of this model is to categorize unstructured text from platforms such as Twitter and YouTube that do not contain information such as the helpfulness reviews or the product ID. Therefore, for the purposes of our model, we will only be using the text content of the reviews as the input for fine-tuning our model and the score for creating the labels.

### 5. Our Pipeline

#### 5.1 Data Preprocessing

The first step in the data pre-processing is to remove all the features that we will not be using for fine-tuning our model.

```
In [5]: # Make a backup of the original dataframe
    df_backup = df.copy()

# List of columns to keep
    columns_to_keep = ['Text', 'Score']

# Drop all other columns
    df_selected = df.drop(columns=[col for col in df.columns if col not in columns)
```

```
# Display the first few rows to verify
print(df_selected.head())
```

```
Score

Text

I have bought several of the Vitality canned d...

Product arrived labeled as Jumbo Salted Peanut...

This is a confection that has been around a fe...

If you are looking for the secret ingredient i...

Great taffy at a great price. There was a wid...
```

Next, we want our model to predict whether a review is negative, neutral, or positive. To do this, we make a score of 1 or 2 represent a review with negative sentiment, a score of 3 represent a review with neutral sentiment, and a review of 4 or 5 to represent a review with positive sentiment.

```
In [6]: import seaborn as sns
        # Function to map 'Score' to 'label' and create a bar chart of label distrib
        def plot sentiment distribution(df, score column='Score', label column='labe
            Maps the score to sentiment labels, creates the 'label' column, and plot
            Parameters:
            df (pd.DataFrame): The input dataset as a pandas DataFrame.
            score column (str): The name of the column containing the score data.
            label column (str): The name of the column to store the sentiment label
            Returns:
            None: Displays a bar chart showing the distribution of sentiment labels.
            # Set the background of the plot to a white grid
            sns.set(style="whitegrid")
            # Define the mapping from Score to label
            score to label = {
                1: 0, # Negative sentiment
                2: 0, # Negative sentiment
                3: 1, # Neutral sentiment
                4: 2, # Positive sentiment
                5: 2 # Positive sentiment
            }
            # Map 'Score' to 'label' column
            df[label column] = df[score column].map(score to label)
            # Count the occurrences of each label
            label counts = df[label column].value counts().sort index()
            # Define label names for better readability
            label names = ['Negative', 'Neutral', 'Positive']
            # Create the bar chart
            plt.figure(figsize=(8, 6))
            bars = plt.bar(label names, label counts, color=['red', 'gray', 'green']
```



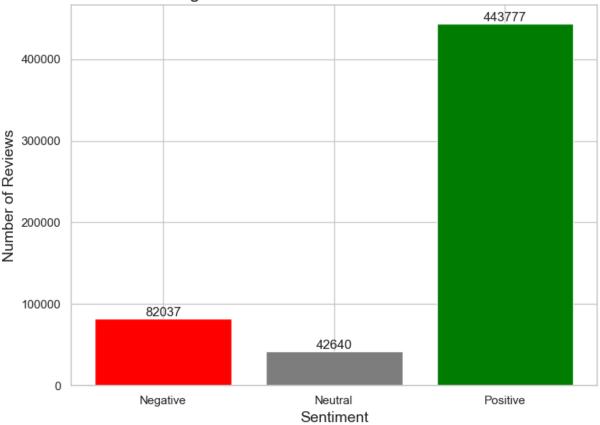


Figure 4 shows the distribution of each label. This plot illustrates the already known issue that we do not have a uniform distribution of the different review sentiments. We fix this by randomly selecting 42,640 samples from the "Positive" sentiment class and the "Negative" sentiment class and dropping the other rows. This fixes two problems for us. The first is the already mentioned non-uniform distribution of the sentiments. If we show our model signficantly more reviews with positive sentiments as opposed to negative or neutral sentiments, the final

model will perform poorly when classifying new reviews with neutral or negative sentiment. The second issue it resolves is it reduces the size of our very large dataset. Due to computational limitations, it would be infeasible to use all the data in the original dataset.

```
In [7]: # Define the target number of samples per sentiment
    target_count = 42640

# Set a random seed for reproducibility
    random_state = 42

# Sample equal number of reviews for each sentiment
    df_balanced = df_selected.groupby('label').apply(
        lambda x: x.sample(n=target_count, random_state=random_state)
).reset_index(drop=True)

# Shuffle the balanced DataFrame
    df_balanced = df_balanced.sample(frac=1, random_state=random_state).reset_ir

# Plot the Sentiment-Distribution
    plot_sentiment_distribution(df_balanced, score_column='Score', label_column=
```

/var/folders/yl/842rvw2j7mq5t40qz5hg61fh0000gn/T/ipykernel\_56877/3387499048.

py:8: DeprecationWarning: DataFrameGroupBy.apply operated on the grouping columns. This behavior is deprecated, and in a future version of pandas the grouping columns will be excluded from the operation. Either pass `include\_groups=False` to exclude the groupings or explicitly select the grouping columns after groupby to silence this warning.

df balanced = df selected.groupby('label').apply(

42640 42640 42640 40000 35000 30000 Number of Reviews 25000 20000 15000 10000 5000 0 Neutral Positive Negative Sentiment

Figure 5: Uniform Distribution of Sentiment Labels

Now that the labels quantities in the dataset have a uniform distribution, the next step is to fix the issue of unwanted characters being included in the text.

```
In [8]:
        import re
        # Define a function to clean text
        def clean text(text):
            # Remove HTML tags
            text = re.sub(r'<.*?>', '', text)
            # Remove leading and trailing whitespace
            text = text.strip()
            # Replace multiple spaces with a single space
            text = re.sub(r'\s+', ' ', text)
            return text
        # Apply the cleaning function to the 'Text' column
        df_balanced['Text'] = df_balanced['Text'].apply(clean_text)
        # Create a word cloud to confirm adjustments of text worked correctly
        text = ' '.join(df_balanced['Text'].astype(str))
        wordcloud = WordCloud(width=800, height=400, background color='white').gener
        # Set the figure size for better readability
        plt.figure(figsize=(12, 6))
        # Plot the wordcloud
        plt.imshow(wordcloud, interpolation='bilinear')
```

```
# Set up axis and labels
plt.axis('off')
plt.title('Figure 6: Word Cloud of Cleaned Reviews', fontsize = 18)
# Show the plot
plt.show()
```



Figue 6 confirms that the desired changes to the dataset have been made. We now define the tokenizer that we will use for our model.

```
In [9]: # Setting up the model and tokenizer
from transformers import TFRobertaForSequenceClassification, RobertaTokenize
# Load the pre-trained tokenizer and sentiment fine-tuned model
tokenizer = RobertaTokenizer.from_pretrained('distilroberta-base')
model = TFRobertaForSequenceClassification.from pretrained('distilroberta-base')
```

/opt/anaconda3/envs/venv\_py311/lib/python3.11/site-packages/transformers/tok enization\_utils\_base.py:1601: FutureWarning: `clean\_up\_tokenization\_spaces` was not set. It will be set to `True` by default. This behavior will be depr acted in transformers v4.45, and will be then set to `False` by default. For more details check this issue: https://github.com/huggingface/transformers/issues/31884

warnings.warn(

All model checkpoint layers were used when initializing TFRobertaForSequence Classification.

All the layers of TFRobertaForSequenceClassification were initialized from the model checkpoint at cardiffnlp/twitter-roberta-base-sentiment. If your task is similar to the task the model of the checkpoint was trained on, you can already use TFRobertaForSequenceClassification for predictions without further training.

The RoBERTa model that we will fine-tune has a maximum input length of 512 tokens. In this next step, we drop any rows that have more than 512 tokens to

account for the fact that the RoBERTa model that we are using takes a maximum of 512 tokens. The results of this can be seen in Figure 6.

```
In [10]: # Function to tokenize text without truncation (to check token length)
def count_tokens(text):
    tokens = tokenizer.encode(text, truncation=False)
    return len(tokens)

# Apply the token counting to the dataset
df_balanced['num_tokens'] = df_balanced['Text'].apply(count_tokens)

# Filter out rows where the token count exceeds max_length
max_length = 512
df_balanced_filtered = df_balanced[df_balanced['num_tokens'] <= max_length].
plot_sentiment_distribution(df_balanced_filtered, score_column='Score', labe</pre>
```

Figure 7: Distribution of Sentiment Labels After Accounting for Text Length

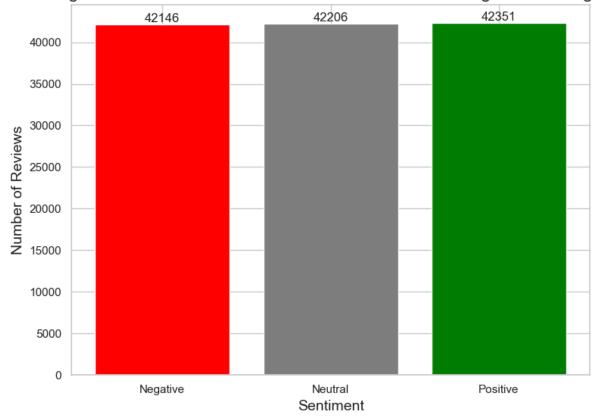
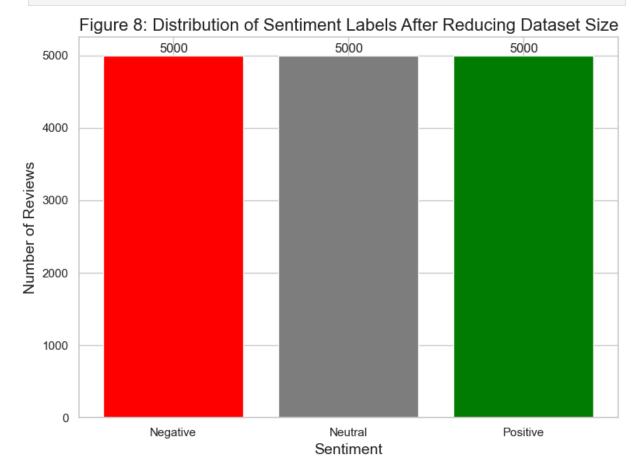


Figure 7 shows that this process had no negative effects on the distribution of sentiments. For the purpose's of training efficiency, we will work with a significantly smaller subset of the data so we can effectively experiemtn with the model. To do this, we will reduce each label to have 5000 rows. Once we have a model that starts seeing good performance, we will scale up the data to try to achieve maximum accuracy.

```
In [45]: # Back up the full sized cleaned dataset
full_cleaned_df = df_balanced_filtered.copy()
```

experimenting\_df = df\_balanced\_filtered.groupby('label').sample(n=5000, ranc
plot\_sentiment\_distribution(experimenting\_df, score\_column='Score', label\_cc



Next, we tokenize the text, split the data into training and testing sets, and convert the resulting data into Tensorflow datasets.

```
In []: # Load the dataset
    df = pd.read_csv('Data/Reviews.csv')

# Map 'Score' to sentiment labels
def map_score_to_label(score):
    if score in [1, 2]:
        return 0 # Negative
    elif score == 3:
        return 1 # Neutral
    else: # 4 or 5
        return 2 # Positive

df['label'] = df['Score'].apply(map_score_to_label)
```

```
In []: # Balance the dataset by sampling 5000 reviews per class
    df_balanced = df.groupby('label').apply(lambda x: x.sample(n=5000, random_st
    # Clean the text
    def clean_text(text):
```

```
# Remove HTML tags
           text = re.sub(r'<.*?>', '', text)
           # Remove non-ASCII characters
           text = text.encode('ascii', 'ignore').decode('utf-8')
           # Remove URLs
           text = re.sub(r'http\S+', '', text)
           # Remove special characters and numbers
           text = re.sub(r'[^A-Za-z\s]', '', text)
           # Convert to lowercase
           text = text.lower()
           # Remove extra spaces
           text = re.sub(r'\s+', ' ', text).strip()
           return text
        df balanced['cleaned text'] = df balanced['Text'].apply(clean text)
In [ ]: # Load the tokenizer
        tokenizer = RobertaTokenizer.from pretrained('distilroberta-base')
        # Tokenize the text
        def tokenize function(texts):
           return tokenizer(
               texts.tolist(),
               padding='max length',
               truncation=True,
               max length=128,
               return tensors='tf'
           )
        tokenized texts = tokenize function(df balanced['cleaned text'])
        # Convert tensors to numpy arrays
        input_ids = tokenized_texts['input_ids'].numpy()
        attention mask = tokenized texts['attention mask'].numpy()
        labels = df balanced['label'].values
        # Clear unused variables to save memory
        del tokenized texts, df balanced
       /usr/local/lib/python3.10/dist-packages/huggingface hub/utils/ token.py:89:
      UserWarning:
      The secret `HF TOKEN` does not exist in your Colab secrets.
      To authenticate with the Hugging Face Hub, create a token in your settings t
      ab (https://huggingface.co/settings/tokens), set it as secret in your Google
      Colab and restart your session.
      You will be able to reuse this secret in all of your notebooks.
      Please note that authentication is recommended but still optional to access
      public models or datasets.
        warnings.warn(
      tokenizer config.json:
                                          | 0.00/25.0 [00:00<?, ?B/s]
                              0%|
      vocab.json: 0%|
      | 0.00/899k [00:00<?, ?B/s]
                                   | 0.00/1.36M [00:00<?, ?B/s]
```

| 0.00/480 [00:00<?, ?B/s]

config.json: 0%|

```
/usr/local/lib/python3.10/dist-packages/transformers/tokenization_utils_bas e.py:1601: FutureWarning: `clean_up_tokenization_spaces` was not set. It wil l be set to `True` by default. This behavior will be depracted in transforme rs v4.45, and will be then set to `False` by default. For more details check this issue: https://github.com/huggingface/transformers/issues/31884 warnings.warn(
```

```
In [ ]: # Split into train, validation, and test sets
        train input ids, temp input ids, train labels, temp labels, train attention
            input ids,
            labels,
            attention mask,
            test size=0.3, # 70% train, 30% temp
            random state=42
        # Further split temp into validation and test sets
        val input ids, test input ids, val labels, test labels, val attention mask,
            temp input ids,
            temp labels,
            temp attention mask,
            test size=0.5, # 15% validation, 15% test
            random state=42
        )
        # Clear unused variables to save memory
        del input ids, attention mask, labels, temp input ids, temp labels, temp att
        # Function to create TensorFlow datasets
        def create tf dataset(input ids, attention mask, labels, batch size):
            dataset = tf.data.Dataset.from tensor slices((
                {'input ids': input ids, 'attention mask': attention mask},
                labels
            ))
            dataset = dataset.batch(batch size)
            return dataset
In [ ]: import tensorflow as tf
        # Build the model function
        def build model(base model, activation function='relu', dropout rate=0.2, us
            # Define inputs
            input ids = tf.keras.Input(shape=(128,), dtype=tf.int32, name='input ids
            attention mask = tf.keras.Input(shape=(128,), dtype=tf.int32, name='atte
            # Get outputs from the base model
            outputs = base model(input ids, attention mask=attention mask)
            cls token = outputs.last hidden state[:, 0, :]
            if additional layers:
              # Add new encoding layers
              x = tf.keras.layers.Dense(512)(cls token)
              if use batchnorm:
                  x = tf.keras.layers.BatchNormalization()(x)
```

x = tf.keras.layers.Activation(activation function)(x)

```
x = tf.keras.layers.Dropout(dropout rate)(x)
              x = tf.keras.layers.Dense(256)(x)
              if use batchnorm:
                  x = tf.keras.layers.BatchNormalization()(x)
              x = tf.keras.layers.Activation(activation function)(x)
              x = tf.keras.layers.Dropout(dropout rate)(x)
              x = tf.keras.layers.Dense(128)(x)
              if use batchnorm:
                  x = tf.keras.layers.BatchNormalization()(x)
              x = tf.keras.layers.Activation(activation function)(x)
              x = tf.keras.layers.Dropout(dropout rate)(x)
              x = tf.keras.layers.Dense(64)(x)
              if use batchnorm:
                  x = tf.keras.layers.BatchNormalization()(x)
              x = tf.keras.layers.Activation(activation function)(x)
              x = tf.keras.layers.Dropout(dropout rate)(x)
              output = tf.keras.layers.Dense(3, activation='softmax')(x)
              output = tf.keras.layers.Dense(3, activation='softmax')(cls token)
            # Build the model
            model = tf.keras.Model(inputs=[input ids, attention mask], outputs=output
            return model
In [ ]: # Function to get optimizer with specified learning rate
        def get optimizer(name, learning rate):
            if name == 'adam':
                return tf.keras.optimizers.Adam(learning rate=learning rate)
            elif name == 'adamw':
                try:
                    # For TensorFlow 2.11 and above
                    return tf.keras.optimizers.experimental.AdamW(learning rate=lear
                except AttributeError:
                    # For earlier versions, use the experimental namespace
                    return tf.keras.optimizers.Adam(learning rate=learning rate) #
            else:
                raise ValueError(f'Unsupported optimizer: {name}')
In [ ]: import random
        from sklearn.metrics import confusion matrix, classification report
        # Define hyperparameter options
        optimizers = ['adam', 'adamw']
        learning rates = [1e-5, 3e-5]
        batch sizes = [16, 32]
        dropout rates = [0.1, 0.2]
        activation functions = ['relu', 'tanh']
        use batchnorm options = [True]
        # Limit the number of configurations to manage computational resources
```

```
max configs = 6 # Adjust this number based on your resources
configs = []
for in range(max configs):
    config = {
        'optimizer': random.choice(optimizers),
        'learning rate': random.choice(learning rates),
        'batch size': random.choice(batch sizes),
        'dropout rate': random.choice(dropout rates),
        'activation function': random.choice(activation functions),
        'use batchnorm': random.choice(use batchnorm options)
    configs.append(config)
# Load the base RoBERTa model and freeze layers
base model = TFRobertaModel.from pretrained('distilroberta-base')
for layer in base model.layers:
    layer.trainable = False
histories = []
results = []
```

Some weights of the PyTorch model were not used when initializing the TF 2.0 model TFRobertaModel: ['lm\_head.bias', 'lm\_head.dense.weight', 'lm\_head.laye r\_norm.bias', 'lm\_head.layer\_norm.weight', 'lm\_head.dense.bias']

- This IS expected if you are initializing TFRobertaModel from a PyTorch mod el trained on another task or with another architecture (e.g. initializing a TFBertForSequenceClassification model from a BertForPreTraining model).

- This IS NOT expected if you are initializing TFRobertaModel from a PyTorch model that you expect to be exactly identical (e.g. initializing a TFBertFor SequenceClassification model from a BertForSequenceClassification model). All the weights of TFRobertaModel were initialized from the PyTorch model. If your task is similar to the task the model of the checkpoint was trained on, you can already use TFRobertaModel for predictions without further train ing.

```
In [ ]: def run model(num epochs, patience, configurations, base model, additional l
          # Initialize variables to keep track of the best model
          best accuracy = 0
          best model = None
          best config = None
          best_y_pred = None
          # Loop over configurations
          for i, config in enumerate(configurations):
              print(f"Training configuration {i+1}/{len(configurations)}: {config}")
              # Build the model
              model = build model(
                  base model,
                  activation function=config['activation function'],
                  dropout rate=config['dropout rate'],
                  use batchnorm=config['use batchnorm'],
                  additional layers= additional layers
              )
```

```
# Get optimizer with specified learning rate
optimizer = get_optimizer(config['optimizer'], learning rate=config['l
# Compile the model
model.compile(
    optimizer=optimizer,
    loss='sparse categorical crossentropy',
   metrics=['accuracy']
)
# Create datasets with specified batch size
batch size = config['batch size']
train dataset = create tf dataset(train input ids, train attention mas
val dataset = create tf dataset(val input ids, val attention mask, val
test dataset = create tf dataset(test input ids, test attention mask,
# Set up callbacks
early stopping = tf.keras.callbacks.EarlyStopping(monitor='val loss',
# Fit the model using the validation set
history = model.fit(
   train dataset,
    epochs=num epochs,
    validation data=val dataset,
    callbacks=[early stopping],
   verbose=1
)
# Record the history and config
histories.append((history, config))
# Evaluate the final model on the test set
loss, accuracy = model.evaluate(test dataset)
print(f"Test accuracy: {accuracy}")
# Store results
results.append({
    'config': config,
    'accuracy': accuracy,
    'loss': loss,
})
# Generate predictions for the test set
y pred probs = model.predict(test dataset)
y pred = np.argmax(y pred probs, axis=1)
# Flatten test labels
y_true = test_labels
# Compute confusion matrix
cm = confusion_matrix(y_true, y_pred)
# Classification report
report = classification_report(y_true, y_pred, target_names=['Negative
print(f"Classification Report for Configuration {i+1}:\n{report}")
```

```
# Plot confusion matrix
   plt.figure(figsize=(6, 5))
   sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=['Negat
   plt.ylabel('Actual')
   plt.xlabel('Predicted')
   plt.title(f'Confusion Matrix for Configuration {i+1}')
   plt.show()
   # Store results
    results.append({
       'config': config,
        'accuracy': accuracy,
        'loss': loss,
        'confusion matrix': cm,
        'classification report': report
   })
   # If current model is better, save it
   if accuracy > best accuracy:
       print("New best model found!")
       # Clear previous best model from memory if exists
       if best_model is not None:
           tf.keras.backend.clear session()
           del best model
       best accuracy = accuracy
       best model = model
       best config = config
       best_y_pred = y pred
       best y true = y true
       # Also store confusion matrix and classification report
       best cm = confusion matrix(y true, y pred)
       best report = classification report(y true, y pred, target names=[
   else:
       # Clear current model from memory
       tf.keras.backend.clear session()
       del model
return best accuracy, best model, best config, best y pred, best y true, b
```

## 6. Training the Model

Below is the code that shows the training process for the first 6 models. After each model is trained, the confusion matrix is printed for the model, along with the combination of hyperparameters used to create it. The first set of 6 models was trained with batch normalization, the second set of 6 models was trained without batch normalization, and the final model was made by freezing only the first 3 encoding layers and unfreezing the 3 closest to the output layer.

#### 6.1 Batch-Normalized Models

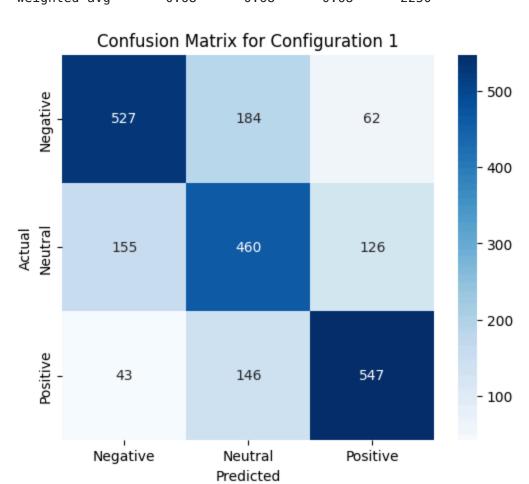
#### 6.1.1 Batch-Normalized Models: Training

```
In [ ]: # Run the model
best_accuracy, best_model, best_config, best_y_pred, best_y_true, best_cm, b
```

```
Training configuration 1/6: {'optimizer': 'adam', 'learning rate': 1e-05, 'b
atch size': 16, 'dropout rate': 0.1, 'activation function': 'relu', 'use bat
chnorm': True}
Epoch 1/80
657/657 [============ ] - 81s 100ms/step - loss: 1.3414 - a
ccuracy: 0.3602 - val loss: 1.0844 - val accuracy: 0.3964
curacy: 0.3923 - val loss: 1.0120 - val accuracy: 0.4631
Epoch 3/80
657/657 [============= ] - 63s 96ms/step - loss: 1.1664 - ac
curacy: 0.4189 - val loss: 0.9703 - val accuracy: 0.5213
Epoch 4/80
curacy: 0.4438 - val loss: 0.9312 - val accuracy: 0.5707
Epoch 5/80
ccuracy: 0.4653 - val loss: 0.8963 - val accuracy: 0.6018
curacy: 0.4993 - val loss: 0.8755 - val accuracy: 0.6129
Epoch 7/80
curacy: 0.5130 - val loss: 0.8504 - val accuracy: 0.6271
Epoch 8/80
curacy: 0.5175 - val loss: 0.8340 - val accuracy: 0.6338
Epoch 9/80
curacy: 0.5286 - val loss: 0.8118 - val accuracy: 0.6471
Epoch 10/80
curacy: 0.5373 - val loss: 0.8015 - val accuracy: 0.6507
Epoch 11/80
curacy: 0.5476 - val loss: 0.7951 - val accuracy: 0.6502
curacy: 0.5554 - val loss: 0.7807 - val accuracy: 0.6551
Epoch 13/80
curacy: 0.5579 - val loss: 0.7744 - val_accuracy: 0.6609
Epoch 14/80
curacy: 0.5650 - val loss: 0.7699 - val accuracy: 0.6640
Epoch 15/80
curacy: 0.5670 - val loss: 0.7594 - val accuracy: 0.6658
Epoch 16/80
657/657 [============] - 63s 96ms/step - loss: 0.9125 - ac
curacy: 0.5749 - val loss: 0.7550 - val accuracy: 0.6640
Epoch 17/80
657/657 [===========] - 63s 97ms/step - loss: 0.9056 - ac
curacy: 0.5763 - val loss: 0.7472 - val accuracy: 0.6720
Epoch 18/80
```

```
curacy: 0.5757 - val loss: 0.7447 - val accuracy: 0.6662
Epoch 19/80
curacy: 0.5884 - val loss: 0.7397 - val accuracy: 0.6676
Epoch 20/80
657/657 [============] - 63s 96ms/step - loss: 0.8941 - ac
curacy: 0.5834 - val_loss: 0.7340 - val accuracy: 0.6707
Epoch 21/80
657/657 [============] - 63s 96ms/step - loss: 0.8808 - ac
curacy: 0.5893 - val loss: 0.7347 - val accuracy: 0.6711
Epoch 22/80
curacy: 0.5928 - val loss: 0.7313 - val accuracy: 0.6716
Epoch 23/80
curacy: 0.5956 - val loss: 0.7275 - val accuracy: 0.6711
curacy: 0.5969 - val loss: 0.7269 - val accuracy: 0.6800
Epoch 25/80
657/657 [============] - 63s 96ms/step - loss: 0.8747 - ac
curacy: 0.5946 - val loss: 0.7220 - val accuracy: 0.6778
Epoch 26/80
curacy: 0.5935 - val loss: 0.7213 - val accuracy: 0.6751
Epoch 27/80
curacy: 0.5981 - val loss: 0.7192 - val accuracy: 0.6751
Epoch 28/80
curacy: 0.6070 - val loss: 0.7145 - val accuracy: 0.6813
657/657 [============== ] - 63s 96ms/step - loss: 0.8500 - ac
curacy: 0.6087 - val loss: 0.7148 - val accuracy: 0.6809
Epoch 30/80
curacy: 0.6201 - val loss: 0.7134 - val accuracy: 0.6778
Epoch 31/80
curacy: 0.6180 - val loss: 0.7050 - val accuracy: 0.6827
Epoch 32/80
curacy: 0.6127 - val loss: 0.7091 - val accuracy: 0.6796
Epoch 33/80
curacy: 0.6180 - val loss: 0.7080 - val accuracy: 0.6804
Epoch 34/80
curacy: 0.6109 - val loss: 0.7084 - val accuracy: 0.6800
Epoch 35/80
curacy: 0.6147 - val loss: 0.7062 - val accuracy: 0.6813
curacy: 0.6818
Test accuracy: 0.6817777752876282
141/141 [============= ] - 12s 71ms/step
```

Classificatio	n Report for	Configur	ation 1:	
	precision	recall	f1-score	support
Negative	0.73	0.68	0.70	773
Neutral	0.58	0.62	0.60	741
Positive	0.74	0.74	0.74	736
accuracy			0.68	2250
macro avg	0.68	0.68	0.68	2250
weighted avg	0.68	0.68	0.68	2250

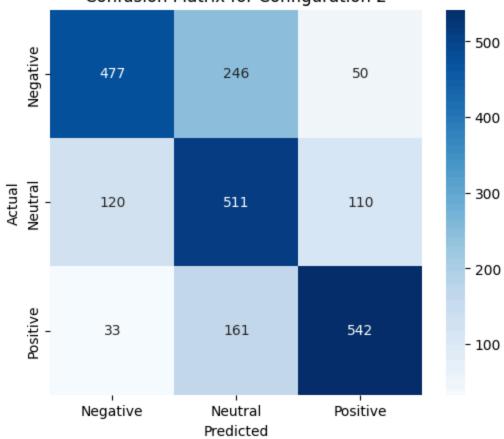


```
New best model found!
Training configuration 2/6: {'optimizer': 'adam', 'learning rate': 1e-05, 'b
atch size': 16, 'dropout rate': 0.2, 'activation function': 'relu', 'use bat
chnorm': True}
Epoch 1/80
657/657 [===========] - 75s 100ms/step - loss: 1.3088 - a
ccuracy: 0.3475 - val loss: 1.0889 - val accuracy: 0.3671
Epoch 2/80
curacy: 0.3684 - val loss: 1.0649 - val accuracy: 0.4169
Epoch 3/80
657/657 [============== ] - 64s 97ms/step - loss: 1.2111 - ac
curacy: 0.3794 - val loss: 1.0413 - val accuracy: 0.4684
Epoch 4/80
curacy: 0.3920 - val loss: 1.0249 - val accuracy: 0.4893
curacy: 0.4123 - val loss: 1.0072 - val accuracy: 0.5036
Epoch 6/80
657/657 [============= ] - 63s 96ms/step - loss: 1.1359 - ac
curacy: 0.4254 - val loss: 0.9859 - val accuracy: 0.5200
Epoch 7/80
curacy: 0.4516 - val loss: 0.9698 - val accuracy: 0.5329
Epoch 8/80
curacy: 0.4690 - val loss: 0.9453 - val accuracy: 0.5551
Epoch 9/80
curacy: 0.4772 - val loss: 0.9302 - val accuracy: 0.5582
curacy: 0.4832 - val loss: 0.9072 - val accuracy: 0.5764
Epoch 11/80
curacy: 0.4989 - val loss: 0.8889 - val accuracy: 0.5880
Epoch 12/80
curacy: 0.5124 - val loss: 0.8775 - val accuracy: 0.5951
Epoch 13/80
curacy: 0.5213 - val loss: 0.8561 - val accuracy: 0.6076
Epoch 14/80
curacy: 0.5190 - val loss: 0.8422 - val accuracy: 0.6164
Epoch 15/80
curacy: 0.5309 - val loss: 0.8300 - val accuracy: 0.6276
Epoch 16/80
curacy: 0.5357 - val loss: 0.8180 - val accuracy: 0.6373
Epoch 17/80
657/657 [===========] - 63s 96ms/step - loss: 0.9613 - ac
curacy: 0.5447 - val loss: 0.8124 - val accuracy: 0.6356
Epoch 18/80
```

```
curacy: 0.5471 - val loss: 0.8090 - val accuracy: 0.6298
Epoch 19/80
657/657 [============= ] - 63s 96ms/step - loss: 0.9607 - ac
curacy: 0.5473 - val loss: 0.7986 - val accuracy: 0.6413
Epoch 20/80
curacy: 0.5464 - val loss: 0.7801 - val accuracy: 0.6480
Epoch 21/80
657/657 [============= ] - 63s 96ms/step - loss: 0.9429 - ac
curacy: 0.5599 - val loss: 0.7774 - val accuracy: 0.6502
curacy: 0.5588 - val loss: 0.7700 - val accuracy: 0.6560
Epoch 23/80
curacy: 0.5685 - val loss: 0.7672 - val accuracy: 0.6578
Epoch 24/80
curacy: 0.5692 - val loss: 0.7654 - val accuracy: 0.6507
Epoch 25/80
curacy: 0.5764 - val loss: 0.7574 - val accuracy: 0.6618
Epoch 26/80
curacy: 0.5740 - val loss: 0.7513 - val accuracy: 0.6693
curacy: 0.5778 - val loss: 0.7499 - val accuracy: 0.6662
Epoch 28/80
curacy: 0.5729 - val loss: 0.7480 - val accuracy: 0.6671
Epoch 29/80
curacy: 0.5823 - val loss: 0.7429 - val accuracy: 0.6693
Epoch 30/80
curacy: 0.5762 - val loss: 0.7327 - val accuracy: 0.6760
Epoch 31/80
657/657 [============] - 63s 96ms/step - loss: 0.9029 - ac
curacy: 0.5825 - val loss: 0.7330 - val accuracy: 0.6751
Epoch 32/80
curacy: 0.5911 - val loss: 0.7333 - val accuracy: 0.6742
Epoch 33/80
657/657 [============= ] - 64s 97ms/step - loss: 0.8874 - ac
curacy: 0.5881 - val loss: 0.7310 - val accuracy: 0.6782
Epoch 34/80
curacy: 0.5951 - val loss: 0.7255 - val accuracy: 0.6751
Epoch 35/80
curacy: 0.5969 - val loss: 0.7253 - val accuracy: 0.6804
Epoch 36/80
curacy: 0.5883 - val loss: 0.7252 - val accuracy: 0.6818
```

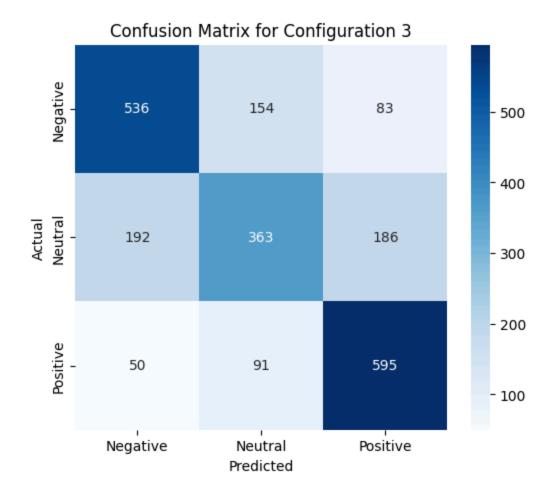
```
Epoch 37/80
657/657 [============= ] - 63s 96ms/step - loss: 0.8825 - ac
curacy: 0.5938 - val loss: 0.7283 - val accuracy: 0.6787
curacy: 0.5997 - val loss: 0.7217 - val accuracy: 0.6840
657/657 [============= ] - 63s 97ms/step - loss: 0.8624 - ac
curacy: 0.6029 - val loss: 0.7148 - val accuracy: 0.6853
Epoch 40/80
657/657 [============] - 63s 96ms/step - loss: 0.8697 - ac
curacy: 0.6002 - val loss: 0.7165 - val accuracy: 0.6836
Epoch 41/80
curacy: 0.6061 - val loss: 0.7144 - val accuracy: 0.6840
Epoch 42/80
curacy: 0.6068 - val loss: 0.7106 - val accuracy: 0.6884
curacy: 0.6056 - val loss: 0.7118 - val accuracy: 0.6853
Epoch 44/80
curacy: 0.6084 - val loss: 0.7070 - val accuracy: 0.6844
Epoch 45/80
curacy: 0.6113 - val loss: 0.7095 - val accuracy: 0.6849
Epoch 46/80
curacy: 0.6120 - val loss: 0.7033 - val accuracy: 0.6871
Epoch 47/80
657/657 [============] - 63s 96ms/step - loss: 0.8490 - ac
curacy: 0.6094 - val loss: 0.7079 - val accuracy: 0.6831
Epoch 48/80
curacy: 0.6132 - val loss: 0.7084 - val accuracy: 0.6871
curacy: 0.6117 - val loss: 0.7050 - val accuracy: 0.6822
Epoch 50/80
curacy: 0.6107 - val loss: 0.7050 - val_accuracy: 0.6818
curacy: 0.6800
Test accuracy: 0.6800000071525574
141/141 [=======] - 12s 72ms/step
Classification Report for Configuration 2:
        precision recall f1-score
                           support
  Negative
           0.76
                 0.62
                       0.68
                              773
           0.56
                 0.69
                       0.62
                              741
  Neutral
  Positive
          0.77
                 0.74
                       0.75
                             736
  accuracy
                       0.68
                             2250
 macro avg 0.70
                 0.68
                       0.68
                             2250
```





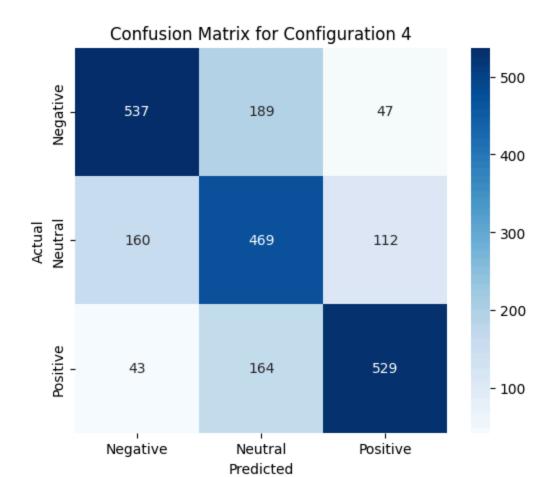
```
Training configuration 3/6: {'optimizer': 'adamw', 'learning_rate': 1e-05,
'batch size': 32, 'dropout rate': 0.1, 'activation function': 'tanh', 'use b
atchnorm': True}
Epoch 1/80
329/329 [============ ] - 71s 186ms/step - loss: 1.1646 - a
ccuracy: 0.4133 - val loss: 0.9168 - val accuracy: 0.5569
329/329 [============ ] - 59s 181ms/step - loss: 1.0241 - a
ccuracy: 0.4986 - val loss: 0.8256 - val accuracy: 0.6178
Epoch 3/80
329/329 [============ ] - 59s 179ms/step - loss: 0.9853 - a
ccuracy: 0.5235 - val loss: 0.7950 - val accuracy: 0.6342
Epoch 4/80
329/329 [=============== ] - 58s 177ms/step - loss: 0.9407 - a
ccuracy: 0.5493 - val loss: 0.7703 - val accuracy: 0.6564
Epoch 5/80
329/329 [============ ] - 58s 177ms/step - loss: 0.9257 - a
ccuracy: 0.5671 - val loss: 0.7540 - val accuracy: 0.6618
329/329 [============= ] - 58s 176ms/step - loss: 0.9061 - a
ccuracy: 0.5725 - val loss: 0.7654 - val accuracy: 0.6609
Epoch 7/80
329/329 [=========== ] - 59s 180ms/step - loss: 0.8978 - a
ccuracy: 0.5814 - val loss: 0.7447 - val accuracy: 0.6640
Epoch 8/80
329/329 [============== ] - 60s 182ms/step - loss: 0.8962 - a
ccuracy: 0.5843 - val loss: 0.7394 - val accuracy: 0.6693
Epoch 9/80
329/329 [============ ] - 59s 178ms/step - loss: 0.8791 - a
ccuracy: 0.5898 - val loss: 0.7354 - val accuracy: 0.6760
Epoch 10/80
329/329 [============= ] - 58s 177ms/step - loss: 0.8792 - a
ccuracy: 0.5903 - val loss: 0.7484 - val accuracy: 0.6684
Epoch 11/80
329/329 [=========== ] - 58s 176ms/step - loss: 0.8602 - a
ccuracy: 0.6021 - val loss: 0.7249 - val accuracy: 0.6818
Epoch 12/80
329/329 [============= ] - 59s 178ms/step - loss: 0.8667 - a
ccuracy: 0.6015 - val loss: 0.7213 - val accuracy: 0.6813
Epoch 13/80
329/329 [============= ] - 58s 177ms/step - loss: 0.8571 - a
ccuracy: 0.6025 - val loss: 0.7157 - val accuracy: 0.6867
Epoch 14/80
329/329 [============ ] - 58s 176ms/step - loss: 0.8540 - a
ccuracy: 0.6099 - val loss: 0.7184 - val accuracy: 0.6791
Epoch 15/80
329/329 [============ ] - 60s 181ms/step - loss: 0.8424 - a
ccuracy: 0.6165 - val loss: 0.7170 - val accuracy: 0.6836
Epoch 16/80
329/329 [============= ] - 58s 176ms/step - loss: 0.8435 - a
ccuracy: 0.6099 - val loss: 0.7229 - val accuracy: 0.6796
Epoch 17/80
329/329 [============= ] - 58s 176ms/step - loss: 0.8460 - a
ccuracy: 0.6117 - val loss: 0.7123 - val accuracy: 0.6844
Epoch 18/80
329/329 [=============== ] - 58s 178ms/step - loss: 0.8430 - a
```

```
ccuracy: 0.6159 - val loss: 0.7031 - val accuracy: 0.6947
Epoch 19/80
329/329 [============= ] - 58s 176ms/step - loss: 0.8379 - a
ccuracy: 0.6222 - val loss: 0.7014 - val accuracy: 0.6933
Epoch 20/80
329/329 [============= ] - 58s 177ms/step - loss: 0.8321 - a
ccuracy: 0.6176 - val loss: 0.7160 - val accuracy: 0.6827
Epoch 21/80
329/329 [=============] - 58s 176ms/step - loss: 0.8240 - a
ccuracy: 0.6237 - val loss: 0.7063 - val accuracy: 0.6907
Epoch 22/80
329/329 [============= ] - 58s 177ms/step - loss: 0.8282 - a
ccuracy: 0.6229 - val loss: 0.6997 - val accuracy: 0.6902
Epoch 23/80
ccuracy: 0.6252 - val loss: 0.6994 - val accuracy: 0.6929
Epoch 24/80
329/329 [============= ] - 58s 177ms/step - loss: 0.8312 - a
ccuracy: 0.6238 - val loss: 0.6977 - val accuracy: 0.6956
Epoch 25/80
329/329 [============ ] - 59s 180ms/step - loss: 0.8231 - a
ccuracy: 0.6315 - val loss: 0.7000 - val accuracy: 0.6911
Epoch 26/80
329/329 [============= ] - 58s 177ms/step - loss: 0.8238 - a
ccuracy: 0.6273 - val loss: 0.7112 - val accuracy: 0.6831
Epoch 27/80
329/329 [=========== ] - 59s 180ms/step - loss: 0.8189 - a
ccuracy: 0.6330 - val_loss: 0.7077 - val_accuracy: 0.6849
Epoch 28/80
329/329 [=========== ] - 58s 176ms/step - loss: 0.8166 - a
ccuracy: 0.6282 - val loss: 0.6983 - val accuracy: 0.6920
71/71 [============ ] - 10s 135ms/step - loss: 0.7444 - acc
uracy: 0.6640
Test accuracy: 0.6639999747276306
71/71 [=======] - 12s 131ms/step
Classification Report for Configuration 3:
           precision recall f1-score
                                       support
   Negative
                0.69
                        0.69
                                 0.69
                                          773
                0.60
                        0.49
                                 0.54
                                          741
    Neutral
   Positive
               0.69
                        0.81
                                 0.74
                                          736
   accuracy
                                 0.66
                                         2250
               0.66
                        0.66
                                 0.66
                                          2250
  macro avg
             0.66
weighted avg
                        0.66
                                 0.66
                                         2250
```



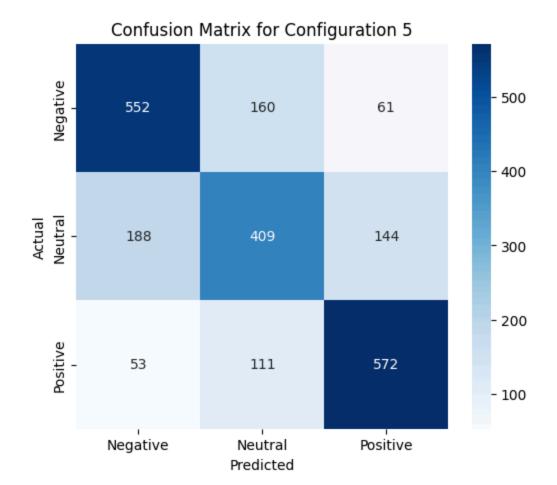
```
Training configuration 4/6: {'optimizer': 'adam', 'learning rate': 3e-05, 'b
atch size': 32, 'dropout rate': 0.1, 'activation function': 'relu', 'use bat
chnorm': True}
Epoch 1/80
329/329 [============= ] - 69s 186ms/step - loss: 1.2325 - a
ccuracy: 0.3845 - val loss: 1.0759 - val accuracy: 0.4009
329/329 [============= ] - 58s 176ms/step - loss: 1.0995 - a
ccuracy: 0.4637 - val loss: 0.9677 - val accuracy: 0.5240
Epoch 3/80
329/329 [============= ] - 58s 176ms/step - loss: 1.0272 - a
ccuracy: 0.5068 - val loss: 0.9125 - val accuracy: 0.5609
Epoch 4/80
329/329 [=============== ] - 58s 177ms/step - loss: 0.9839 - a
ccuracy: 0.5324 - val loss: 0.8582 - val accuracy: 0.5884
Epoch 5/80
329/329 [============ ] - 59s 180ms/step - loss: 0.9542 - a
ccuracy: 0.5526 - val loss: 0.8174 - val accuracy: 0.6218
329/329 [============= ] - 59s 178ms/step - loss: 0.9307 - a
ccuracy: 0.5630 - val loss: 0.7895 - val accuracy: 0.6364
Epoch 7/80
329/329 [=========== ] - 58s 176ms/step - loss: 0.9017 - a
ccuracy: 0.5827 - val loss: 0.7576 - val accuracy: 0.6511
Epoch 8/80
329/329 [=============== ] - 59s 178ms/step - loss: 0.8905 - a
ccuracy: 0.5927 - val loss: 0.7424 - val accuracy: 0.6622
Epoch 9/80
329/329 [============= ] - 58s 177ms/step - loss: 0.8758 - a
ccuracy: 0.5964 - val loss: 0.7455 - val accuracy: 0.6609
Epoch 10/80
329/329 [============= ] - 58s 176ms/step - loss: 0.8596 - a
ccuracy: 0.6075 - val loss: 0.7302 - val accuracy: 0.6676
Epoch 11/80
329/329 [=========== ] - 59s 178ms/step - loss: 0.8585 - a
ccuracy: 0.6048 - val loss: 0.7128 - val accuracy: 0.6782
Epoch 12/80
329/329 [============== ] - 60s 181ms/step - loss: 0.8542 - a
ccuracy: 0.6094 - val loss: 0.7047 - val accuracy: 0.6889
Epoch 13/80
329/329 [============= ] - 59s 181ms/step - loss: 0.8429 - a
ccuracy: 0.6165 - val loss: 0.7019 - val accuracy: 0.6893
Epoch 14/80
329/329 [============= ] - 58s 176ms/step - loss: 0.8335 - a
ccuracy: 0.6172 - val loss: 0.6947 - val accuracy: 0.6831
Epoch 15/80
329/329 [============ ] - 58s 176ms/step - loss: 0.8286 - a
ccuracy: 0.6240 - val loss: 0.6930 - val accuracy: 0.6831
Epoch 16/80
329/329 [============= ] - 58s 176ms/step - loss: 0.8269 - a
ccuracy: 0.6218 - val loss: 0.6910 - val accuracy: 0.6884
Epoch 17/80
329/329 [============= ] - 58s 177ms/step - loss: 0.8261 - a
ccuracy: 0.6211 - val_loss: 0.6891 - val_accuracy: 0.6804
Epoch 18/80
329/329 [================== ] - 59s 178ms/step - loss: 0.8168 - a
```

```
ccuracy: 0.6271 - val loss: 0.6865 - val accuracy: 0.6849
Epoch 19/80
329/329 [============= ] - 58s 177ms/step - loss: 0.8082 - a
ccuracy: 0.6328 - val loss: 0.6838 - val accuracy: 0.6924
Epoch 20/80
329/329 [============= ] - 58s 176ms/step - loss: 0.8109 - a
ccuracy: 0.6311 - val loss: 0.6822 - val accuracy: 0.6893
Epoch 21/80
329/329 [============= ] - 58s 176ms/step - loss: 0.7962 - a
ccuracy: 0.6431 - val loss: 0.6836 - val accuracy: 0.6947
Epoch 22/80
329/329 [============= ] - 59s 180ms/step - loss: 0.8017 - a
ccuracy: 0.6379 - val loss: 0.6822 - val accuracy: 0.6924
Epoch 23/80
ccuracy: 0.6390 - val loss: 0.6798 - val accuracy: 0.6938
Epoch 24/80
329/329 [============= ] - 58s 176ms/step - loss: 0.7858 - a
ccuracy: 0.6442 - val loss: 0.6751 - val accuracy: 0.6947
Epoch 25/80
329/329 [============= ] - 58s 176ms/step - loss: 0.7825 - a
ccuracy: 0.6501 - val loss: 0.6792 - val accuracy: 0.6902
Epoch 26/80
329/329 [============= ] - 58s 177ms/step - loss: 0.7823 - a
ccuracy: 0.6487 - val loss: 0.6706 - val accuracy: 0.6978
Epoch 27/80
329/329 [============ ] - 59s 180ms/step - loss: 0.7838 - a
ccuracy: 0.6482 - val loss: 0.6847 - val accuracy: 0.6876
Epoch 28/80
329/329 [============ ] - 59s 180ms/step - loss: 0.7813 - a
ccuracy: 0.6476 - val loss: 0.6747 - val accuracy: 0.6938
Epoch 29/80
329/329 [============= ] - 58s 175ms/step - loss: 0.7813 - a
ccuracy: 0.6487 - val loss: 0.6785 - val accuracy: 0.6964
Epoch 30/80
329/329 [============= ] - 58s 176ms/step - loss: 0.7787 - a
ccuracy: 0.6518 - val loss: 0.6717 - val accuracy: 0.6956
uracy: 0.6822
Test accuracy: 0.6822222471237183
71/71 [========= ] - 24s 131ms/step
Classification Report for Configuration 4:
           precision recall f1-score
                                     support
   Negative
               0.73
                       0.69
                                0.71
                                         773
                                0.60
   Neutral
               0.57
                       0.63
                                         741
   Positive
               0.77
                                        736
                       0.72
                                0.74
                                0.68
                                        2250
   accuracy
  macro avg
              0.69
                       0.68
                                0.68
                                        2250
weighted avg 0.69
                       0.68
                              0.68
                                        2250
```



```
New best model found!
Training configuration 5/6: {'optimizer': 'adamw', 'learning rate': 1e-05,
'batch size': 16, 'dropout rate': 0.1, 'activation function': 'tanh', 'use b
atchnorm': True}
Epoch 1/80
657/657 [===========] - 74s 100ms/step - loss: 1.1607 - a
ccuracy: 0.4116 - val loss: 0.8530 - val accuracy: 0.6124
Epoch 2/80
curacy: 0.5026 - val loss: 0.7762 - val accuracy: 0.6493
Epoch 3/80
657/657 [============= ] - 63s 96ms/step - loss: 0.9790 - ac
curacy: 0.5250 - val loss: 0.7458 - val accuracy: 0.6569
Epoch 4/80
curacy: 0.5524 - val loss: 0.7274 - val accuracy: 0.6693
657/657 [============] - 63s 97ms/step - loss: 0.9222 - ac
curacy: 0.5597 - val loss: 0.7226 - val accuracy: 0.6769
Epoch 6/80
657/657 [============] - 64s 97ms/step - loss: 0.9090 - ac
curacy: 0.5733 - val loss: 0.7170 - val accuracy: 0.6733
Epoch 7/80
curacy: 0.5810 - val loss: 0.7155 - val accuracy: 0.6716
Epoch 8/80
curacy: 0.5907 - val loss: 0.7105 - val accuracy: 0.6778
Epoch 9/80
curacy: 0.5890 - val loss: 0.7080 - val accuracy: 0.6849
curacy: 0.5950 - val loss: 0.6973 - val accuracy: 0.6898
Epoch 11/80
curacy: 0.5947 - val loss: 0.6973 - val accuracy: 0.6947
Epoch 12/80
curacy: 0.6015 - val loss: 0.6896 - val accuracy: 0.6902
Epoch 13/80
curacy: 0.6028 - val loss: 0.6916 - val accuracy: 0.6907
Epoch 14/80
curacy: 0.6082 - val loss: 0.6890 - val accuracy: 0.6920
Epoch 15/80
657/657 [============== ] - 64s 97ms/step - loss: 0.8550 - ac
curacy: 0.6141 - val loss: 0.6916 - val accuracy: 0.6902
Epoch 16/80
curacy: 0.6125 - val loss: 0.6894 - val accuracy: 0.6902
Epoch 17/80
657/657 [============] - 64s 97ms/step - loss: 0.8397 - ac
curacy: 0.6161 - val loss: 0.6891 - val accuracy: 0.6911
Epoch 18/80
```

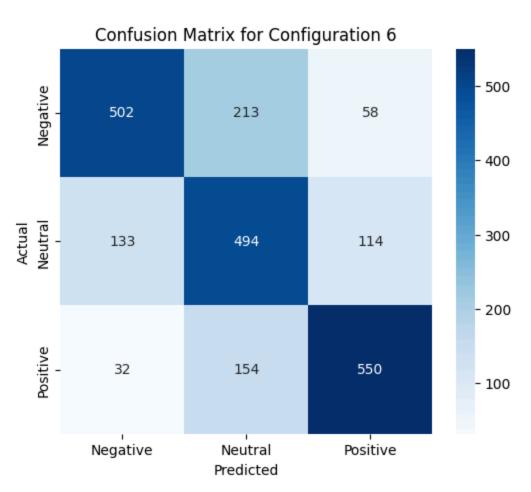
```
curacy: 0.6118 - val loss: 0.6872 - val accuracy: 0.6916
Epoch 19/80
curacy: 0.6201 - val loss: 0.6838 - val accuracy: 0.6938
Epoch 20/80
curacy: 0.6230 - val loss: 0.6844 - val accuracy: 0.6978
Epoch 21/80
curacy: 0.6188 - val loss: 0.6932 - val accuracy: 0.6898
curacy: 0.6245 - val loss: 0.6777 - val accuracy: 0.6982
Epoch 23/80
curacy: 0.6294 - val loss: 0.6829 - val accuracy: 0.6947
Epoch 24/80
curacy: 0.6194 - val loss: 0.6795 - val accuracy: 0.6973
Epoch 25/80
curacy: 0.6230 - val loss: 0.6801 - val accuracy: 0.6978
Epoch 26/80
curacy: 0.6337 - val loss: 0.6783 - val accuracy: 0.6996
curacy: 0.6813
Test accuracy: 0.6813333630561829
141/141 [========== ] - 12s 71ms/step
Classification Report for Configuration 5:
       precision recall f1-score
                        support
          0.70
               0.71
                     0.70
  Negative
                          773
  Neutral
          0.60
               0.55
                     0.58
                          741
  Positive
          0.74
               0.78
                    0.76
                          736
                    0.68
                          2250
  accuracy
 macro avq
          0.68
               0.68
                    0.68
                          2250
          0.68
                    0.68
weighted avg
               0.68
                          2250
```



```
Training configuration 6/6: {'optimizer': 'adam', 'learning rate': 3e-05, 'b
atch size': 16, 'dropout rate': 0.2, 'activation function': 'relu', 'use bat
chnorm': True}
Epoch 1/80
657/657 [============ ] - 74s 100ms/step - loss: 1.3482 - a
ccuracy: 0.3550 - val loss: 1.0332 - val accuracy: 0.4782
curacy: 0.4230 - val loss: 0.9304 - val accuracy: 0.5756
Epoch 3/80
curacy: 0.4662 - val loss: 0.8691 - val accuracy: 0.6076
Epoch 4/80
curacy: 0.5013 - val loss: 0.8196 - val accuracy: 0.6378
Epoch 5/80
curacy: 0.5194 - val loss: 0.7807 - val accuracy: 0.6600
curacy: 0.5428 - val loss: 0.7592 - val accuracy: 0.6702
Epoch 7/80
curacy: 0.5544 - val loss: 0.7450 - val accuracy: 0.6698
Epoch 8/80
curacy: 0.5575 - val loss: 0.7329 - val accuracy: 0.6653
Epoch 9/80
curacy: 0.5627 - val loss: 0.7243 - val accuracy: 0.6773
Epoch 10/80
curacy: 0.5732 - val loss: 0.7209 - val accuracy: 0.6764
Epoch 11/80
curacy: 0.5789 - val loss: 0.7185 - val accuracy: 0.6773
curacy: 0.5835 - val loss: 0.7123 - val accuracy: 0.6764
Epoch 13/80
curacy: 0.5838 - val loss: 0.7122 - val accuracy: 0.6796
Epoch 14/80
curacy: 0.5937 - val loss: 0.7048 - val accuracy: 0.6831
Epoch 15/80
curacy: 0.5977 - val loss: 0.7074 - val accuracy: 0.6867
Epoch 16/80
curacy: 0.6047 - val loss: 0.6997 - val accuracy: 0.6867
Epoch 17/80
curacy: 0.6018 - val loss: 0.6946 - val accuracy: 0.6871
Epoch 18/80
```

```
curacy: 0.6086 - val loss: 0.6944 - val accuracy: 0.6849
Epoch 19/80
curacy: 0.6124 - val loss: 0.6930 - val accuracy: 0.6876
Epoch 20/80
657/657 [============] - 64s 98ms/step - loss: 0.8491 - ac
curacy: 0.6149 - val_loss: 0.6916 - val accuracy: 0.6867
Epoch 21/80
curacy: 0.6109 - val loss: 0.6912 - val accuracy: 0.6920
Epoch 22/80
curacy: 0.6280 - val loss: 0.6951 - val accuracy: 0.6858
Epoch 23/80
curacy: 0.6178 - val loss: 0.6902 - val accuracy: 0.6964
curacy: 0.6224 - val loss: 0.6880 - val accuracy: 0.6933
Epoch 25/80
curacy: 0.6270 - val loss: 0.6832 - val accuracy: 0.7009
Epoch 26/80
curacy: 0.6305 - val loss: 0.6822 - val accuracy: 0.6911
Epoch 27/80
curacy: 0.6237 - val loss: 0.6784 - val accuracy: 0.6978
Epoch 28/80
curacy: 0.6400 - val loss: 0.6772 - val accuracy: 0.6973
curacy: 0.6313 - val loss: 0.6852 - val accuracy: 0.6996
Epoch 30/80
curacy: 0.6349 - val loss: 0.6762 - val accuracy: 0.6982
Epoch 31/80
curacy: 0.6309 - val loss: 0.6787 - val accuracy: 0.6951
Epoch 32/80
curacy: 0.6282 - val loss: 0.6730 - val accuracy: 0.6987
Epoch 33/80
curacy: 0.6426 - val loss: 0.6755 - val accuracy: 0.7022
Epoch 34/80
657/657 [=============] - 63s 97ms/step - loss: 0.7926 - ac
curacy: 0.6450 - val loss: 0.6740 - val accuracy: 0.6964
Epoch 35/80
curacy: 0.6419 - val loss: 0.6704 - val accuracy: 0.6987
Epoch 36/80
657/657 [============] - 63s 96ms/step - loss: 0.7974 - ac
curacy: 0.6413 - val loss: 0.6775 - val accuracy: 0.6996
Epoch 37/80
```

```
curacy: 0.6467 - val loss: 0.6720 - val accuracy: 0.7044
Epoch 38/80
curacy: 0.6415 - val_loss: 0.6745 - val_accuracy: 0.7031
Epoch 39/80
curacy: 0.6440 - val_loss: 0.6706 - val_accuracy: 0.7013
curacy: 0.6871
Test accuracy: 0.6871111392974854
Classification Report for Configuration 6:
        precision
               recall f1-score
                      0.70
  Negative
          0.75
                0.65
                            773
  Neutral
          0.57
                0.67
                      0.62
                            741
  Positive
          0.76
                0.75
                      0.75
                            736
  accuracy
                      0.69
                            2250
          0.70
                0.69
                      0.69
 macro avg
                            2250
weighted avg
          0.70
                0.69
                      0.69
                            2250
```



New best model found!

6.1.2 Batch-Normalized Models: Examples of Mis-classified Points

```
In [ ]: # After the training loop
        print(f"\nBest Model Configuration: {best config}")
        print(f"Best Model Test Accuracy: {best accuracy}")
        # Use the best model's predictions
        y pred = best y pred
        y true = best y true
        # Identify misclassified examples
        misclassified indices = np.where(y pred != y true)[0]
        # Decode test texts
        def decode texts(input ids):
            return [tokenizer.decode(ids, skip special tokens=True) for ids in input
        test texts = decode texts(test input ids)
        # Extract false positives and false negatives
        false positives = []
        false negatives = []
        # Mapping of label indices to label names
        label map = {0: 'Negative', 1: 'Neutral', 2: 'Positive'}
        # Loop through misclassified examples to separate false positives and false
        for idx in misclassified indices:
            true label = y true[idx]
            predicted label = y pred[idx]
            text = test texts[idx]
            if predicted label == 2 and true label != 2:
                # Model predicted Positive, but true label is Negative or Neutral
                false positives append((text, label map[true label], label map[predi
            elif predicted label != 2 and true label == 2:
                # Model predicted Negative or Neutral, but true label is Positive
                false negatives append((text, label map[true label], label map[predi
        # Display a few examples of false positives
        print("\nExamples of False Positives:")
        for i in range(min(3, len(false positives))):
            text, true label, predicted label = false positives[i]
            print(f"\nText: {text}")
            print(f"True Label: {true label}")
            print(f"Predicted Label: {predicted label}")
        # Display a few examples of false negatives
        print("\nExamples of False Negatives:")
        for i in range(min(3, len(false negatives))):
            text, true label, predicted label = false negatives[i]
            print(f"\nText: {text}")
            print(f"True Label: {true label}")
            print(f"Predicted Label: {predicted label}")
```

Best Model Configuration: {'optimizer': 'adam', 'learning\_rate': 3e-05, 'bat ch\_size': 16, 'dropout\_rate': 0.2, 'activation\_function': 'relu', 'use\_batch norm': True}

Best Model Test Accuracy: 0.6871111392974854

Examples of False Positives:

Text: this is quite good ive taken to starting my mornings with a nice hot  ${\tt m}$ 

ug

True Label: Neutral

Predicted Label: Positive

Text: i love gloria jeans cinnamon nut strudel coffee and decided to try this the flavoring is a little over powering for me if youve tried this flavor before and like it you wont be disappointed with it in the kcups packaging

True Label: Neutral Predicted Label: Positive

Text: we love clif kid z bars i decided to try these based on all of the gre at reviews we bought the strawberry and no one in the house likes them my ki ds are picky so i always try the treats just to see if they are even worth t rying i eat healthy and dont mind most healthy snacks but even i did not car e for these i am going to try another flavor but the strawberry just didnt c ut it here

True Label: Negative Predicted Label: Positive

Examples of False Negatives:

Text: i really only needed the funnelpitcher but it was easy and neat to mak e funnel cakes i had been using a kitchen funnel which made a big mess no matter how careful i was i didnt use the ring because i made the funnel cakes in a small deep fryer the mix was fine but i prefer my own recipe if i could have just bought the funnelpitcher that would have been fine and saved me so me money but im not disappointed

True Label: Positive Predicted Label: Neutral

Text: i was very impressed with the flavor growing up on mrs butterworths i developed a taste for commercial syrup and had a hard time adjusting to the real thing but when i opened up the jug of coombs it was delish i would high ly recommend this product

True Label: Positive Predicted Label: Neutral

Text: i got a larger one than my puppy probably required and she loves dragg ing it around playing tug of war and just chewing on it unfortunately it does shed i dont know if there are rope toys out there that are any sturdier but this one does leave little threads behind otherwise great toy

True Label: Positive Predicted Label: Neutral

```
In [ ]: # Save the best model to a .keras file
   best_model.save('best_model.keras')
   print("\nBest model saved to 'best_model.keras'.")
```

```
/usr/local/lib/python3.10/dist-packages/transformers/generation/tf_utils.py:
465: UserWarning: `seed_generator` is deprecated and will be removed in a future version.

warnings.warn("`seed_generator` is deprecated and will be removed in a future version.", UserWarning)

Best model saved to 'best model.keras'.
```

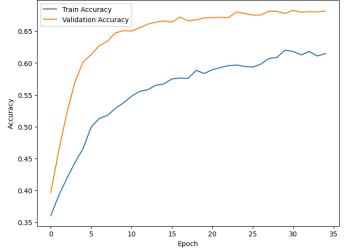
#### 6.1.3 Batch-Normalized Models: Visualizations

We will show the process of training and plotting the training (not exactly sure what she means by this since the next section is to plot the validation set and analyze the behaviour of training and validation, but we'll figure it out).

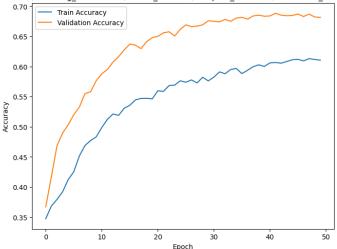
```
# Visualize the results

# Plot training and validation accuracy for each configuration
for i, (history, config) in enumerate(histories):
    plt.figure(figsize=(8, 6))
    plt.plot(history.history['accuracy'], label='Train Accuracy')
    plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
    plt.title(f"Configuration {i+1}: {config}")
    plt.xlabel('Epoch')
    plt.ylabel('Accuracy')
    plt.legend()
    plt.show()
```

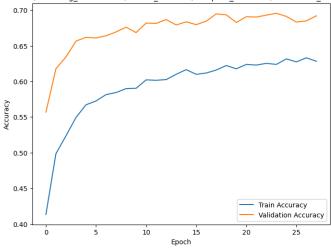
Configuration 1: {'optimizer': 'adam', 'learning\_rate': 1e-05, 'batch\_size': 16, 'dropout\_rate': 0.1, 'activation\_function': 'relu', 'use\_batchnorm': True}



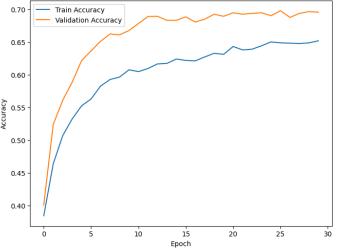
Configuration 2: {'optimizer': 'adam', 'learning\_rate': 1e-05, 'batch\_size': 16, 'dropout\_rate': 0.2, 'activation\_function': 'relu', 'use\_batchnorm': True}



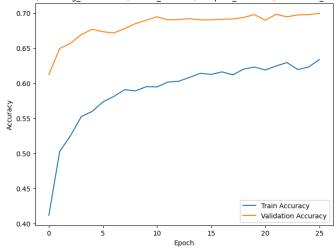
Configuration 3: {'optimizer': 'adamw', 'learning\_rate': 1e-05, 'batch\_size': 32, 'dropout\_rate': 0.1, 'activation\_function': 'tanh', 'use\_batchnorm': True}



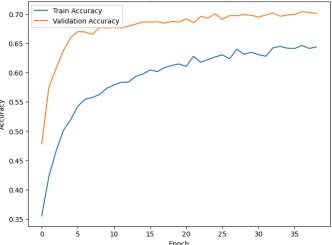
Configuration 4: {'optimizer': 'adam', 'learning\_rate': 3e-05, 'batch\_size': 32, 'dropout\_rate': 0.1, 'activation\_function': 'relu', 'use\_batchnorm': True}



Configuration 5: {'optimizer': 'adamw', 'learning rate': 1e-05, 'batch size': 16, 'dropout rate': 0.1, 'activation function': 'tanh', 'use batchnorm': True}

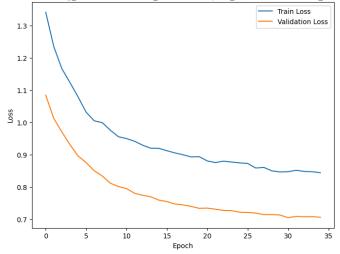


Configuration 6: {'optimizer': 'adam', 'learning\_rate': 3e-05, 'batch\_size': 16, 'dropout\_rate': 0.2, 'activation\_function': 'relu', 'use\_batchnorm': True}

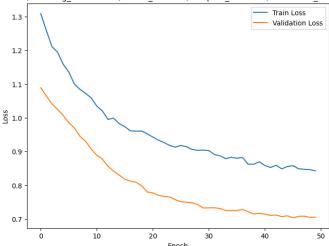


```
In []: # Plot training and validation loss for each configuration
    for i, (history, config) in enumerate(histories):
        plt.figure(figsize=(8, 6))
        plt.plot(history.history['loss'], label='Train Loss')
        plt.plot(history.history['val_loss'], label='Validation Loss')
        plt.title(f"Configuration {i+1}: {config}")
        plt.xlabel('Epoch')
        plt.ylabel('Loss')
        plt.legend()
        plt.show()
```

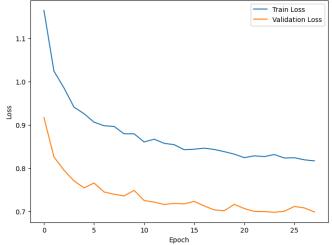
Configuration 1: {'optimizer': 'adam', 'learning\_rate': 1e-05, 'batch\_size': 16, 'dropout\_rate': 0.1, 'activation\_function': 'relu', 'use\_batchnorm': True}



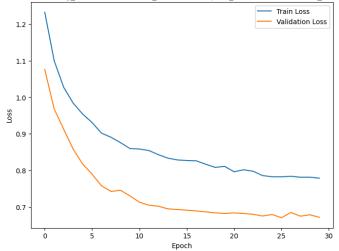
Configuration 2: {'optimizer': 'adam', 'learning\_rate': 1e-05, 'batch\_size': 16, 'dropout\_rate': 0.2, 'activation\_function': 'relu', 'use\_batchnorm': True}



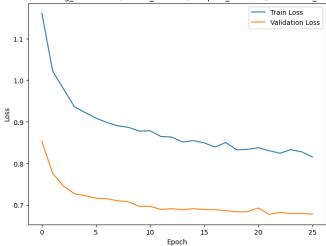
Configuration 3: {'optimizer': 'adamw', 'learning\_rate': 1e-05, 'batch\_size': 32, 'dropout\_rate': 0.1, 'activation\_function': 'tanh', 'use\_batchnorm': True}



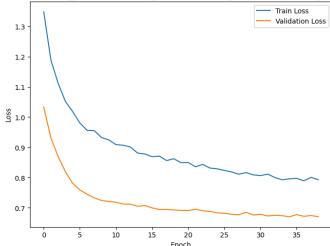
Configuration 4: {'optimizer': 'adam', 'learning\_rate': 3e-05, 'batch\_size': 32, 'dropout\_rate': 0.1, 'activation\_function': 'relu', 'use\_batchnorm': True}



Configuration 5: {'optimizer': 'adamw', 'learning\_rate': 1e-05, 'batch\_size': 16, 'dropout\_rate': 0.1, 'activation\_function': 'tanh', 'use\_batchnorm': True}



Configuration 6: {'optimizer': 'adam', 'learning\_rate': 3e-05, 'batch\_size': 16, 'dropout\_rate': 0.2, 'activation\_function': 'relu', 'use\_batchnorm': True}



### 6.1.4 Batch-Normalized Models: Table of Models

```
In []: # Create a DataFrame of the results
    results_df = pd.DataFrame([{
        'optimizer': res['config']['optimizer'],
        'learning_rate': res['config']['learning_rate'],
        'batch_size': res['config']['batch_size'],
```

```
'dropout_rate': res['config']['dropout_rate'],
     'activation function': res['config']['activation function'],
     'use batchnorm': res['config']['use batchnorm'],
     'accuracy': res['accuracy'],
     'loss': res['loss']
 } for res in results])
 print(results df)
   optimizer learning rate batch size dropout rate activation function \
0
                    0.00001
                                     16
                                                  0.1
        adam
                                                                     relu
                                                  0.1
1
        adam
                    0.00001
                                     16
                                                                     relu
2
        adam
                    0.00001
                                     16
                                                  0.2
                                                                     relu
3
                                     16
                                                  0.2
        adam
                    0.00001
                                                                     relu
                                     32
                                                  0.1
4
       adamw
                    0.00001
                                                                     tanh
5
       adamw
                                     32
                    0.00001
                                                  0.1
                                                                     tanh
                                     32
6
                                                  0.1
        adam
                    0.00003
                                                                     relu
7
        adam
                    0.00003
                                     32
                                                  0.1
                                                                     relu
8
                                                  0.1
       adamw
                    0.00001
                                     16
                                                                     tanh
9
       adamw
                    0.00001
                                     16
                                                  0.1
                                                                     tanh
10
                                     16
                                                  0.2
        adam
                    0.00003
                                                                     relu
                                     16
                                                  0.2
11
        adam
                    0.00003
                                                                     relu
    use batchnorm accuracy
                                 loss
0
             True 0.681778 0.726639
1
             True 0.681778 0.726639
             True 0.680000 0.732409
2
3
             True 0.680000 0.732409
4
             True 0.664000 0.744429
5
            True 0.664000 0.744429
            True 0.682222 0.708122
6
             True 0.682222 0.708122
7
8
            True 0.681333 0.724130
             True 0.681333 0.724130
9
            True 0.687111 0.702287
10
11
            True 0.687111 0.702287
```

### 6.2 Non-Batch-Normalized Models

### 6.2.1 Non-Batch-Normalized Models: Training

```
In [ ]: # Define hyperparameter options
        optimizers = ['adam', 'adamw']
        learning rates = [1e-5, 3e-5]
        batch sizes = [16, 32]
        dropout rates = [0.1, 0.2]
        activation functions = ['relu', 'tanh']
        use batchnorm options = [False] # prevent underfitting
        # Limit the number of configurations to manage computational resources
        max configs = 6 # Adjust this number based on your resources
        configs = []
        for in range(max configs):
            config = {
                'optimizer': random.choice(optimizers),
                'learning rate': random.choice(learning rates),
                'batch size': random.choice(batch sizes),
                 'dropout rate': random.choice(dropout rates),
                 'activation function': random.choice(activation functions),
                'use batchnorm': random.choice(use batchnorm options)
            configs.append(config)
        # Load the base RoBERTa model and freeze layers
        base model = TFRobertaModel.from pretrained('distilroberta-base')
        for layer in base model.layers:
            layer.trainable = False
        histories = []
        results = []
```

Some weights of the PyTorch model were not used when initializing the TF 2.0 model TFRobertaModel: ['lm\_head.bias', 'lm\_head.dense.weight', 'lm\_head.laye r\_norm.bias', 'lm\_head.layer\_norm.weight', 'lm\_head.dense.bias']

- This IS expected if you are initializing TFRobertaModel from a PyTorch mod el trained on another task or with another architecture (e.g. initializing a TFBertForSequenceClassification model from a BertForPreTraining model).

- This IS NOT expected if you are initializing TFRobertaModel from a PyTorch model that you expect to be exactly identical (e.g. initializing a TFBertFor SequenceClassification model from a BertForSequenceClassification model). All the weights of TFRobertaModel were initialized from the PyTorch model. If your task is similar to the task the model of the checkpoint was trained on, you can already use TFRobertaModel for predictions without further train ing.

```
In []: def run_model(num_epochs, patience, configurations, base_model, additional_l
    # Initialize variables to keep track of the best model
    best_accuracy = 0
    best_model = None
    best_config = None
    best_y_pred = None

# Loop over configurations
for i, config in enumerate(configurations):
```

```
print(f"Training configuration {i+1}/{len(configurations)}: {config}")
# Build the model
model = build model(
   base model,
    activation function=config['activation function'],
    dropout rate=config['dropout rate'],
    use batchnorm=config['use batchnorm'],
    additional layers = additional layers
)
# Get optimizer with specified learning rate
optimizer = get optimizer(config['optimizer'], learning rate=config['l
# Compile the model
model.compile(
    optimizer=optimizer,
   loss='sparse categorical crossentropy',
   metrics=['accuracy']
)
# Create datasets with specified batch size
batch size = config['batch size']
train dataset = create tf dataset(train input ids, train attention mas
val_dataset = create_tf_dataset(val_input_ids, val_attention_mask, val
test dataset = create tf dataset(test input ids, test attention mask,
# Set up callbacks
early stopping = tf.keras.callbacks.EarlyStopping(monitor='val loss',
# Fit the model using the validation set
history = model.fit(
   train dataset,
    epochs=num epochs,
   validation data=val dataset,
    callbacks=[early stopping],
   verbose=1
)
# Record the history and config
histories.append((history, config))
# Evaluate the final model on the test set
loss, accuracy = model.evaluate(test dataset)
print(f"Test accuracy: {accuracy}")
# Store results
results.append({
    'config': config,
    'accuracy': accuracy,
    'loss': loss,
})
# Generate predictions for the test set
y pred probs = model.predict(test dataset)
y pred = np.argmax(y pred probs, axis=1)
```

```
# Flatten test labels
   y true = test labels
   # Compute confusion matrix
   cm = confusion matrix(y true, y pred)
   # Classification report
    report = classification report(y true, y pred, target names=['Negative
   print(f"Classification Report for Configuration {i+1}:\n{report}")
   # Plot confusion matrix
   plt.figure(figsize=(6, 5))
   sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=['Negat']
   plt.ylabel('Actual')
   plt.xlabel('Predicted')
   plt.title(f'Confusion Matrix for Configuration {i+1}')
   plt.show()
   # Store results
    results.append({
        'config': config,
        'accuracy': accuracy,
        'loss': loss,
        'confusion matrix': cm,
        'classification report': report
   })
   # If current model is better, save it
   if accuracy > best accuracy:
       print("New best model found!")
       # Clear previous best model from memory if exists
       if best model is not None:
            tf.keras.backend.clear session()
            del best model
       best accuracy = accuracy
       best model = model
       best config = config
       best y pred = y pred
       best y true = y true
       # Also store confusion matrix and classification report
       best cm = confusion matrix(y true, y pred)
       best report = classification report(y true, y pred, target names=[
   else:
       # Clear current model from memory
       tf.keras.backend.clear session()
       del model
return best accuracy, best model, best config, best y pred, best y true, b
```

In [ ]: # Run the model
 best\_accuracy, best\_model, best\_config, best\_y\_pred, best\_y\_true, best\_cm, b

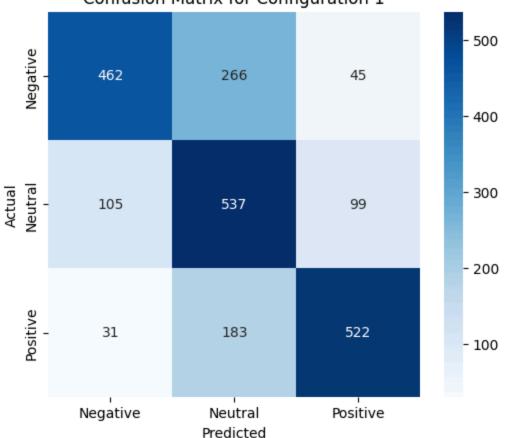
```
Training configuration 1/6: {'optimizer': 'adamw', 'learning_rate': 3e-05,
'batch size': 16, 'dropout rate': 0.1, 'activation function': 'tanh', 'use b
atchnorm': False}
Epoch 1/80
657/657 [===========] - 73s 101ms/step - loss: 1.0993 - a
ccuracy: 0.3907 - val loss: 0.9956 - val accuracy: 0.5182
657/657 [============== ] - 65s 99ms/step - loss: 0.9773 - ac
curacy: 0.5142 - val loss: 0.8534 - val accuracy: 0.6067
Epoch 3/80
curacy: 0.5789 - val loss: 0.7991 - val accuracy: 0.6120
Epoch 4/80
curacy: 0.6171 - val loss: 0.7604 - val accuracy: 0.6422
Epoch 5/80
curacy: 0.6250 - val loss: 0.7507 - val accuracy: 0.6476
curacy: 0.6375 - val loss: 0.7394 - val accuracy: 0.6542
Epoch 7/80
curacy: 0.6441 - val loss: 0.7198 - val accuracy: 0.6671
Epoch 8/80
curacy: 0.6417 - val loss: 0.7224 - val accuracy: 0.6698
Epoch 9/80
curacy: 0.6493 - val loss: 0.7193 - val accuracy: 0.6662
Epoch 10/80
curacy: 0.6508 - val loss: 0.7222 - val accuracy: 0.6680
Epoch 11/80
curacy: 0.6549 - val loss: 0.7002 - val accuracy: 0.6853
curacy: 0.6570 - val loss: 0.7019 - val accuracy: 0.6849
Epoch 13/80
curacy: 0.6632 - val loss: 0.7052 - val accuracy: 0.6773
Epoch 14/80
curacy: 0.6637 - val loss: 0.6993 - val accuracy: 0.6853
Epoch 15/80
657/657 [============] - 65s 99ms/step - loss: 0.7620 - ac
curacy: 0.6630 - val loss: 0.6938 - val accuracy: 0.6867
Epoch 16/80
657/657 [============] - 66s 100ms/step - loss: 0.7544 - a
ccuracy: 0.6654 - val loss: 0.6934 - val accuracy: 0.6844
Epoch 17/80
657/657 [============] - 65s 98ms/step - loss: 0.7601 - ac
curacy: 0.6652 - val loss: 0.6895 - val accuracy: 0.6933
Epoch 18/80
```

curacy: 0.6631 - val loss: 0.6907 - val accuracy: 0.6889 Epoch 19/80 curacy: 0.6692 - val loss: 0.6926 - val accuracy: 0.6867 Epoch 20/80 657/657 [============] - 64s 97ms/step - loss: 0.7508 - ac curacy: 0.6688 - val loss: 0.6899 - val accuracy: 0.6880 Epoch 21/80 curacy: 0.6662 - val loss: 0.6903 - val accuracy: 0.6907 curacy: 0.6760 Test accuracy: 0.6759999990463257 141/141 [========= ] - 12s 73ms/step

Classification Report for Configuration 1:

	precision	recall	f1-score	support
Negative Neutral Positive	0.77 0.54 0.78	0.60 0.72 0.71	0.67 0.62 0.74	773 741 736
accuracy macro avg weighted avg	0.70 0.70	0.68 0.68	0.68 0.68 0.68	2250 2250 2250

## Confusion Matrix for Configuration 1



```
New best model found!
Training configuration 2/6: {'optimizer': 'adamw', 'learning rate': 3e-05,
'batch size': 32, 'dropout rate': 0.1, 'activation function': 'tanh', 'use b
atchnorm': False}
Epoch 1/80
329/329 [============= ] - 72s 194ms/step - loss: 1.1209 - a
ccuracy: 0.3676 - val loss: 1.0282 - val accuracy: 0.5738
Epoch 2/80
329/329 [============ ] - 60s 184ms/step - loss: 1.0380 - a
ccuracy: 0.4625 - val loss: 0.9283 - val accuracy: 0.5907
Epoch 3/80
329/329 [============= ] - 61s 186ms/step - loss: 0.9422 - a
ccuracy: 0.5429 - val loss: 0.8324 - val accuracy: 0.6298
Epoch 4/80
329/329 [============ ] - 60s 182ms/step - loss: 0.8771 - a
ccuracy: 0.5938 - val loss: 0.7770 - val accuracy: 0.6453
Epoch 5/80
329/329 [============= ] - 61s 184ms/step - loss: 0.8437 - a
ccuracy: 0.6037 - val loss: 0.7508 - val accuracy: 0.6631
Epoch 6/80
329/329 [============= ] - 60s 183ms/step - loss: 0.8198 - a
ccuracy: 0.6279 - val loss: 0.7322 - val accuracy: 0.6764
Epoch 7/80
329/329 [============= ] - 61s 186ms/step - loss: 0.8069 - a
ccuracy: 0.6335 - val loss: 0.7208 - val accuracy: 0.6796
Epoch 8/80
329/329 [============= ] - 61s 185ms/step - loss: 0.8039 - a
ccuracy: 0.6338 - val loss: 0.7129 - val accuracy: 0.6813
Epoch 9/80
329/329 [============= ] - 60s 183ms/step - loss: 0.7923 - a
ccuracy: 0.6405 - val loss: 0.7175 - val accuracy: 0.6778
329/329 [============== ] - 60s 183ms/step - loss: 0.7834 - a
ccuracy: 0.6409 - val loss: 0.7073 - val accuracy: 0.6884
Epoch 11/80
ccuracy: 0.6470 - val loss: 0.7088 - val accuracy: 0.6849
Epoch 12/80
329/329 [============ ] - 60s 183ms/step - loss: 0.7698 - a
ccuracy: 0.6571 - val loss: 0.7033 - val accuracy: 0.6880
Epoch 13/80
329/329 [================== ] - 62s 187ms/step - loss: 0.7778 - a
ccuracy: 0.6505 - val loss: 0.7005 - val accuracy: 0.6911
Epoch 14/80
329/329 [============== ] - 60s 183ms/step - loss: 0.7690 - a
ccuracy: 0.6579 - val loss: 0.6930 - val accuracy: 0.6951
Epoch 15/80
329/329 [=============== ] - 61s 184ms/step - loss: 0.7641 - a
ccuracy: 0.6619 - val loss: 0.6933 - val accuracy: 0.6933
Epoch 16/80
329/329 [============ ] - 61s 185ms/step - loss: 0.7692 - a
ccuracy: 0.6568 - val loss: 0.6938 - val accuracy: 0.6933
Epoch 17/80
329/329 [============= ] - 61s 185ms/step - loss: 0.7666 - a
ccuracy: 0.6602 - val loss: 0.6976 - val accuracy: 0.6933
Epoch 18/80
```

```
329/329 [============== ] - 60s 183ms/step - loss: 0.7639 - a
ccuracy: 0.6630 - val loss: 0.6894 - val accuracy: 0.6924
Epoch 19/80
329/329 [============= ] - 60s 184ms/step - loss: 0.7599 - a
ccuracy: 0.6647 - val loss: 0.6859 - val accuracy: 0.6978
Epoch 20/80
329/329 [================= ] - 61s 185ms/step - loss: 0.7608 - a
ccuracy: 0.6610 - val loss: 0.6957 - val accuracy: 0.6898
Epoch 21/80
329/329 [============= ] - 60s 182ms/step - loss: 0.7588 - a
ccuracy: 0.6638 - val loss: 0.6916 - val accuracy: 0.6960
ccuracy: 0.6661 - val loss: 0.6838 - val accuracy: 0.7040
Epoch 23/80
329/329 [============= ] - 60s 183ms/step - loss: 0.7509 - a
ccuracy: 0.6687 - val loss: 0.6808 - val accuracy: 0.7000
Epoch 24/80
ccuracy: 0.6690 - val loss: 0.6809 - val accuracy: 0.7031
Epoch 25/80
329/329 [============ ] - 60s 183ms/step - loss: 0.7475 - a
ccuracy: 0.6700 - val loss: 0.6879 - val accuracy: 0.6960
Epoch 26/80
329/329 [============= ] - 60s 183ms/step - loss: 0.7523 - a
ccuracy: 0.6660 - val loss: 0.6822 - val accuracy: 0.7000
Epoch 27/80
329/329 [============= ] - 60s 183ms/step - loss: 0.7454 - a
ccuracy: 0.6707 - val_loss: 0.6770 - val_accuracy: 0.7071
Epoch 28/80
ccuracy: 0.6703 - val loss: 0.6865 - val accuracy: 0.6978
Epoch 29/80
329/329 [=========== ] - 61s 185ms/step - loss: 0.7455 - a
ccuracy: 0.6693 - val loss: 0.6821 - val accuracy: 0.6978
Epoch 30/80
329/329 [============== ] - 61s 184ms/step - loss: 0.7472 - a
ccuracy: 0.6678 - val loss: 0.6844 - val accuracy: 0.6969
Epoch 31/80
329/329 [============= ] - 62s 188ms/step - loss: 0.7436 - a
ccuracy: 0.6716 - val loss: 0.6756 - val accuracy: 0.7027
Epoch 32/80
329/329 [============== ] - 60s 182ms/step - loss: 0.7443 - a
ccuracy: 0.6681 - val loss: 0.6736 - val accuracy: 0.7036
Epoch 33/80
329/329 [============= ] - 61s 185ms/step - loss: 0.7423 - a
ccuracy: 0.6772 - val loss: 0.6753 - val accuracy: 0.7018
Epoch 34/80
329/329 [============= ] - 61s 187ms/step - loss: 0.7494 - a
ccuracy: 0.6681 - val loss: 0.6767 - val accuracy: 0.7000
Epoch 35/80
329/329 [============ ] - 60s 182ms/step - loss: 0.7412 - a
ccuracy: 0.6720 - val loss: 0.6802 - val accuracy: 0.6947
Epoch 36/80
329/329 [============== ] - 60s 184ms/step - loss: 0.7406 - a
ccuracy: 0.6704 - val loss: 0.6726 - val accuracy: 0.7027
```

```
Epoch 37/80
329/329 [============ ] - 60s 183ms/step - loss: 0.7415 - a
ccuracy: 0.6709 - val loss: 0.6751 - val accuracy: 0.7013
Epoch 38/80
329/329 [============ ] - 61s 185ms/step - loss: 0.7388 - a
ccuracy: 0.6723 - val loss: 0.6742 - val accuracy: 0.6982
329/329 [============= ] - 60s 182ms/step - loss: 0.7421 - a
ccuracy: 0.6699 - val loss: 0.6779 - val accuracy: 0.6960
Epoch 40/80
329/329 [============ ] - 60s 183ms/step - loss: 0.7393 - a
ccuracy: 0.6774 - val loss: 0.6749 - val accuracy: 0.7013
71/71 [============ ] - 10s 138ms/step - loss: 0.7090 - acc
uracy: 0.6844
Test accuracy: 0.6844444274902344
Classification Report for Configuration 2:
                       recall f1-score
            precision
                                       support
   Negative
                0.76
                         0.63
                                 0.69
                                           773
                0.57
                                 0.62
    Neutral
                         0.67
                                           741
   Positive
                0.76
                         0.76
                                 0.76
                                           736
   accuracy
                                 0.68
                                          2250
  macro avq
                0.69
                         0.69
                                 0.69
                                          2250
```

0.69

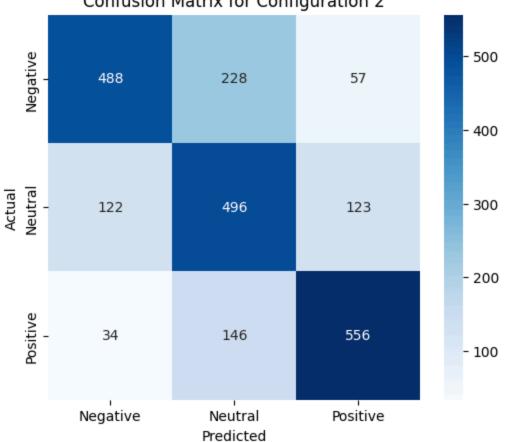
2250



0.68

0.70

weighted avg

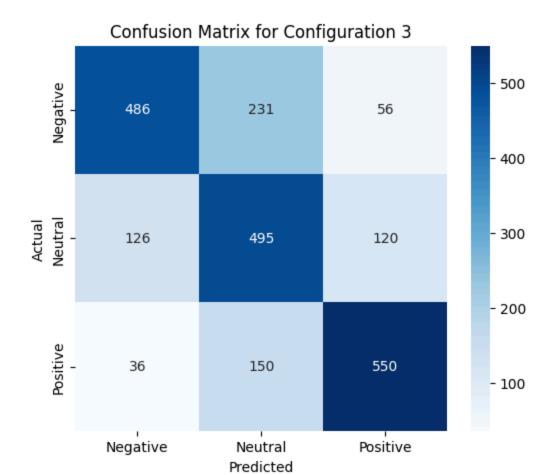


```
New best model found!
Training configuration 3/6: {'optimizer': 'adamw', 'learning rate': 1e-05,
'batch size': 32, 'dropout rate': 0.1, 'activation function': 'tanh', 'use b
atchnorm': False}
Epoch 1/80
329/329 [============= ] - 73s 193ms/step - loss: 1.1361 - a
ccuracy: 0.3505 - val loss: 1.0738 - val accuracy: 0.4000
Epoch 2/80
329/329 [============ ] - 62s 187ms/step - loss: 1.1099 - a
ccuracy: 0.3741 - val loss: 1.0474 - val accuracy: 0.5320
Epoch 3/80
329/329 [============= ] - 60s 183ms/step - loss: 1.0849 - a
ccuracy: 0.4051 - val loss: 1.0149 - val accuracy: 0.5960
Epoch 4/80
329/329 [============ ] - 60s 183ms/step - loss: 1.0496 - a
ccuracy: 0.4467 - val loss: 0.9792 - val accuracy: 0.6044
Epoch 5/80
329/329 [============= ] - 60s 184ms/step - loss: 1.0173 - a
ccuracy: 0.4775 - val loss: 0.9390 - val accuracy: 0.6227
Epoch 6/80
329/329 [============= ] - 60s 184ms/step - loss: 0.9832 - a
ccuracy: 0.5094 - val loss: 0.9004 - val accuracy: 0.6302
Epoch 7/80
329/329 [============= ] - 61s 185ms/step - loss: 0.9536 - a
ccuracy: 0.5308 - val loss: 0.8660 - val accuracy: 0.6396
Epoch 8/80
329/329 [============ ] - 61s 186ms/step - loss: 0.9247 - a
ccuracy: 0.5549 - val loss: 0.8363 - val accuracy: 0.6422
Epoch 9/80
329/329 [============ ] - 61s 186ms/step - loss: 0.9010 - a
ccuracy: 0.5650 - val loss: 0.8120 - val accuracy: 0.6462
329/329 [=============== ] - 61s 186ms/step - loss: 0.8818 - a
ccuracy: 0.5771 - val loss: 0.7985 - val accuracy: 0.6462
Epoch 11/80
329/329 [============== ] - 61s 186ms/step - loss: 0.8631 - a
ccuracy: 0.5956 - val loss: 0.7805 - val accuracy: 0.6533
Epoch 12/80
329/329 [============ ] - 60s 183ms/step - loss: 0.8477 - a
ccuracy: 0.6029 - val loss: 0.7685 - val accuracy: 0.6560
Epoch 13/80
329/329 [=================== ] - 61s 186ms/step - loss: 0.8468 - a
ccuracy: 0.6103 - val loss: 0.7598 - val accuracy: 0.6622
Epoch 14/80
329/329 [=================== ] - 61s 186ms/step - loss: 0.8340 - a
ccuracy: 0.6140 - val loss: 0.7503 - val accuracy: 0.6649
Epoch 15/80
329/329 [============= ] - 60s 183ms/step - loss: 0.8351 - a
ccuracy: 0.6130 - val loss: 0.7435 - val accuracy: 0.6689
Epoch 16/80
329/329 [============= ] - 60s 183ms/step - loss: 0.8267 - a
ccuracy: 0.6209 - val loss: 0.7404 - val accuracy: 0.6684
Epoch 17/80
329/329 [============ ] - 60s 183ms/step - loss: 0.8197 - a
ccuracy: 0.6273 - val loss: 0.7397 - val accuracy: 0.6707
Epoch 18/80
```

```
ccuracy: 0.6286 - val loss: 0.7303 - val accuracy: 0.6773
Epoch 19/80
329/329 [============] - 60s 183ms/step - loss: 0.8092 - a
ccuracy: 0.6284 - val loss: 0.7276 - val accuracy: 0.6760
Epoch 20/80
329/329 [================== ] - 61s 186ms/step - loss: 0.8087 - a
ccuracy: 0.6252 - val loss: 0.7249 - val accuracy: 0.6751
Epoch 21/80
329/329 [============= ] - 61s 185ms/step - loss: 0.8006 - a
ccuracy: 0.6372 - val loss: 0.7235 - val accuracy: 0.6782
ccuracy: 0.6372 - val_loss: 0.7184 - val_accuracy: 0.6809
Epoch 23/80
329/329 [============= ] - 62s 187ms/step - loss: 0.7977 - a
ccuracy: 0.6368 - val loss: 0.7168 - val accuracy: 0.6791
Epoch 24/80
ccuracy: 0.6385 - val loss: 0.7120 - val accuracy: 0.6827
Epoch 25/80
329/329 [============= ] - 60s 183ms/step - loss: 0.7844 - a
ccuracy: 0.6480 - val loss: 0.7113 - val accuracy: 0.6836
Epoch 26/80
329/329 [============= ] - 60s 184ms/step - loss: 0.7907 - a
ccuracy: 0.6415 - val loss: 0.7086 - val accuracy: 0.6840
Epoch 27/80
329/329 [============= ] - 60s 183ms/step - loss: 0.7879 - a
ccuracy: 0.6443 - val_loss: 0.7094 - val_accuracy: 0.6822
Epoch 28/80
329/329 [============ ] - 60s 183ms/step - loss: 0.7863 - a
ccuracy: 0.6454 - val loss: 0.7066 - val accuracy: 0.6849
Epoch 29/80
329/329 [============ ] - 61s 186ms/step - loss: 0.7842 - a
ccuracy: 0.6460 - val loss: 0.7091 - val accuracy: 0.6849
Epoch 30/80
329/329 [============= ] - 60s 183ms/step - loss: 0.7775 - a
ccuracy: 0.6551 - val loss: 0.7024 - val accuracy: 0.6880
Epoch 31/80
329/329 [============ ] - 60s 182ms/step - loss: 0.7888 - a
ccuracy: 0.6456 - val loss: 0.7029 - val accuracy: 0.6840
Epoch 32/80
329/329 [============== ] - 60s 183ms/step - loss: 0.7811 - a
ccuracy: 0.6513 - val loss: 0.7036 - val accuracy: 0.6867
Epoch 33/80
329/329 [============ ] - 60s 183ms/step - loss: 0.7752 - a
ccuracy: 0.6521 - val_loss: 0.6987 - val_accuracy: 0.6884
Epoch 34/80
329/329 [============= ] - 61s 185ms/step - loss: 0.7759 - a
ccuracy: 0.6562 - val loss: 0.7007 - val accuracy: 0.6889
Epoch 35/80
329/329 [============ ] - 60s 183ms/step - loss: 0.7756 - a
ccuracy: 0.6577 - val loss: 0.7008 - val accuracy: 0.6889
Epoch 36/80
329/329 [============== ] - 61s 186ms/step - loss: 0.7736 - a
ccuracy: 0.6626 - val loss: 0.6966 - val accuracy: 0.6902
```

```
Epoch 37/80
329/329 [============ ] - 60s 183ms/step - loss: 0.7680 - a
ccuracy: 0.6582 - val loss: 0.6961 - val accuracy: 0.6907
Epoch 38/80
329/329 [============= ] - 60s 184ms/step - loss: 0.7687 - a
ccuracy: 0.6596 - val loss: 0.6949 - val accuracy: 0.6920
329/329 [============= ] - 60s 183ms/step - loss: 0.7742 - a
ccuracy: 0.6533 - val loss: 0.6936 - val accuracy: 0.6933
Epoch 40/80
329/329 [============= ] - 60s 183ms/step - loss: 0.7707 - a
ccuracy: 0.6583 - val loss: 0.6950 - val accuracy: 0.6951
Epoch 41/80
329/329 [=============== ] - 60s 184ms/step - loss: 0.7708 - a
ccuracy: 0.6521 - val loss: 0.6932 - val accuracy: 0.6920
Epoch 42/80
329/329 [============= ] - 60s 183ms/step - loss: 0.7722 - a
ccuracy: 0.6561 - val loss: 0.6937 - val accuracy: 0.6938
329/329 [============= ] - 61s 186ms/step - loss: 0.7690 - a
ccuracy: 0.6592 - val loss: 0.6924 - val accuracy: 0.6964
Epoch 44/80
ccuracy: 0.6608 - val loss: 0.6942 - val accuracy: 0.6942
Epoch 45/80
329/329 [============== ] - 61s 184ms/step - loss: 0.7627 - a
ccuracy: 0.6572 - val loss: 0.6927 - val accuracy: 0.6933
Epoch 46/80
ccuracy: 0.6599 - val loss: 0.6896 - val accuracy: 0.6960
Epoch 47/80
329/329 [============= ] - 60s 182ms/step - loss: 0.7632 - a
ccuracy: 0.6594 - val loss: 0.6922 - val accuracy: 0.6938
Epoch 48/80
329/329 [============ ] - 61s 187ms/step - loss: 0.7617 - a
ccuracy: 0.6627 - val loss: 0.6901 - val accuracy: 0.6960
329/329 [================= ] - 61s 185ms/step - loss: 0.7638 - a
ccuracy: 0.6605 - val loss: 0.6877 - val accuracy: 0.6960
Epoch 50/80
329/329 [============= ] - 61s 186ms/step - loss: 0.7611 - a
ccuracy: 0.6611 - val loss: 0.6879 - val accuracy: 0.6973
Epoch 51/80
329/329 [============= ] - 61s 186ms/step - loss: 0.7566 - a
ccuracy: 0.6668 - val loss: 0.6865 - val accuracy: 0.6938
Epoch 52/80
329/329 [============= ] - 61s 186ms/step - loss: 0.7543 - a
ccuracy: 0.6674 - val loss: 0.6868 - val accuracy: 0.6951
Epoch 53/80
329/329 [================== ] - 61s 185ms/step - loss: 0.7556 - a
ccuracy: 0.6629 - val loss: 0.6886 - val accuracy: 0.6982
Epoch 54/80
329/329 [============== ] - 60s 183ms/step - loss: 0.7520 - a
ccuracy: 0.6677 - val loss: 0.6861 - val accuracy: 0.6969
329/329 [============== ] - 61s 186ms/step - loss: 0.7600 - a
```

```
ccuracy: 0.6619 - val loss: 0.6853 - val accuracy: 0.6964
Epoch 56/80
329/329 [============= ] - 61s 185ms/step - loss: 0.7530 - a
ccuracy: 0.6676 - val loss: 0.6831 - val accuracy: 0.6947
Epoch 57/80
329/329 [============= ] - 61s 187ms/step - loss: 0.7577 - a
ccuracy: 0.6670 - val loss: 0.6821 - val accuracy: 0.6969
Epoch 58/80
329/329 [============= ] - 61s 185ms/step - loss: 0.7543 - a
ccuracy: 0.6716 - val loss: 0.6836 - val accuracy: 0.6960
Epoch 59/80
329/329 [============ ] - 61s 184ms/step - loss: 0.7498 - a
ccuracy: 0.6684 - val loss: 0.6843 - val accuracy: 0.6987
Epoch 60/80
329/329 [=========== ] - 60s 182ms/step - loss: 0.7582 - a
ccuracy: 0.6682 - val loss: 0.6834 - val accuracy: 0.6973
Epoch 61/80
329/329 [=========== ] - 61s 185ms/step - loss: 0.7504 - a
ccuracy: 0.6683 - val loss: 0.6828 - val accuracy: 0.6978
71/71 [============= ] - 10s 140ms/step - loss: 0.7195 - acc
uracy: 0.6804
Test accuracy: 0.6804444193840027
71/71 [========] - 11s 136ms/step
Classification Report for Configuration 3:
            precision recall f1-score
                                        support
                      0.63
0.67
                0.75
                                 0.68
                                           773
   Negative
    Neutral
                0.57
                                 0.61
                                           741
   Positive
                0.76
                        0.75
                                 0.75
                                          736
                                 0.68
                                          2250
   accuracy
              0.69
                                 0.68
  macro avq
                       0.68
                                          2250
weighted avg
               0.69
                         0.68
                                 0.68
                                          2250
```



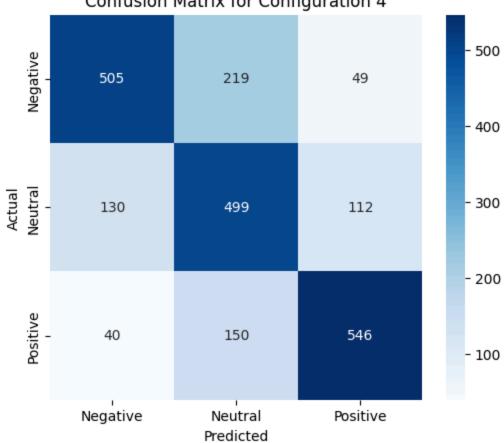
```
Training configuration 4/6: {'optimizer': 'adam', 'learning rate': 3e-05, 'b
atch size': 32, 'dropout rate': 0.2, 'activation function': 'relu', 'use bat
chnorm': False}
Epoch 1/80
329/329 [============= ] - 69s 191ms/step - loss: 1.1155 - a
ccuracy: 0.3318 - val loss: 1.0964 - val accuracy: 0.3342
329/329 [============ ] - 60s 183ms/step - loss: 1.1025 - a
ccuracy: 0.3490 - val loss: 1.0937 - val accuracy: 0.3324
Epoch 3/80
329/329 [============= ] - 62s 187ms/step - loss: 1.0987 - a
ccuracy: 0.3498 - val loss: 1.0883 - val accuracy: 0.3840
Epoch 4/80
329/329 [============== ] - 60s 183ms/step - loss: 1.0907 - a
ccuracy: 0.3711 - val loss: 1.0680 - val accuracy: 0.4951
Epoch 5/80
329/329 [============= ] - 61s 186ms/step - loss: 1.0743 - a
ccuracy: 0.4113 - val loss: 1.0329 - val accuracy: 0.5276
329/329 [============ ] - 60s 183ms/step - loss: 1.0440 - a
ccuracy: 0.4538 - val loss: 0.9725 - val accuracy: 0.5800
Epoch 7/80
329/329 [============ ] - 61s 184ms/step - loss: 1.0057 - a
ccuracy: 0.4925 - val loss: 0.9113 - val accuracy: 0.5809
Epoch 8/80
329/329 [=============== ] - 62s 187ms/step - loss: 0.9552 - a
ccuracy: 0.5353 - val loss: 0.8570 - val accuracy: 0.6058
Epoch 9/80
329/329 [============ ] - 61s 185ms/step - loss: 0.9235 - a
ccuracy: 0.5483 - val loss: 0.8266 - val accuracy: 0.6249
Epoch 10/80
329/329 [=============] - 61s 187ms/step - loss: 0.9071 - a
ccuracy: 0.5683 - val loss: 0.8044 - val accuracy: 0.6311
Epoch 11/80
329/329 [=========== ] - 60s 183ms/step - loss: 0.8834 - a
ccuracy: 0.5806 - val loss: 0.7896 - val accuracy: 0.6382
Epoch 12/80
329/329 [============== ] - 61s 184ms/step - loss: 0.8752 - a
ccuracy: 0.5900 - val loss: 0.7806 - val accuracy: 0.6436
Epoch 13/80
329/329 [=========== ] - 61s 184ms/step - loss: 0.8619 - a
ccuracy: 0.5988 - val loss: 0.7704 - val accuracy: 0.6480
Epoch 14/80
329/329 [============ ] - 60s 181ms/step - loss: 0.8536 - a
ccuracy: 0.6059 - val loss: 0.7618 - val accuracy: 0.6569
Epoch 15/80
329/329 [============ ] - 60s 183ms/step - loss: 0.8500 - a
ccuracy: 0.6100 - val loss: 0.7476 - val accuracy: 0.6622
Epoch 16/80
329/329 [============] - 59s 180ms/step - loss: 0.8396 - a
ccuracy: 0.6161 - val_loss: 0.7480 - val accuracy: 0.6676
Epoch 17/80
329/329 [============= ] - 60s 184ms/step - loss: 0.8270 - a
ccuracy: 0.6250 - val_loss: 0.7364 - val_accuracy: 0.6733
Epoch 18/80
329/329 [============== ] - 61s 185ms/step - loss: 0.8216 - a
```

```
ccuracy: 0.6270 - val loss: 0.7375 - val accuracy: 0.6742
Epoch 19/80
329/329 [============ ] - 60s 183ms/step - loss: 0.8101 - a
ccuracy: 0.6372 - val loss: 0.7276 - val accuracy: 0.6733
Epoch 20/80
329/329 [============ ] - 60s 183ms/step - loss: 0.8089 - a
ccuracy: 0.6321 - val loss: 0.7247 - val accuracy: 0.6720
Epoch 21/80
329/329 [============= ] - 61s 186ms/step - loss: 0.8082 - a
ccuracy: 0.6360 - val loss: 0.7208 - val accuracy: 0.6822
Epoch 22/80
329/329 [============= ] - 60s 181ms/step - loss: 0.8000 - a
ccuracy: 0.6409 - val loss: 0.7230 - val accuracy: 0.6773
Epoch 23/80
329/329 [============= ] - 60s 183ms/step - loss: 0.7949 - a
ccuracy: 0.6411 - val loss: 0.7104 - val accuracy: 0.6884
Epoch 24/80
329/329 [============ ] - 60s 182ms/step - loss: 0.8006 - a
ccuracy: 0.6474 - val loss: 0.7097 - val accuracy: 0.6880
Epoch 25/80
329/329 [============ ] - 60s 182ms/step - loss: 0.7916 - a
ccuracy: 0.6463 - val loss: 0.7085 - val accuracy: 0.6831
Epoch 26/80
329/329 [============= ] - 61s 185ms/step - loss: 0.7864 - a
ccuracy: 0.6458 - val loss: 0.7031 - val accuracy: 0.6836
Epoch 27/80
329/329 [============ ] - 59s 181ms/step - loss: 0.7835 - a
ccuracy: 0.6528 - val loss: 0.7034 - val accuracy: 0.6822
Epoch 28/80
329/329 [============= ] - 60s 183ms/step - loss: 0.7830 - a
ccuracy: 0.6551 - val loss: 0.7039 - val accuracy: 0.6844
329/329 [============== ] - 59s 180ms/step - loss: 0.7778 - a
ccuracy: 0.6572 - val loss: 0.6976 - val accuracy: 0.6813
Epoch 30/80
329/329 [============== ] - 59s 180ms/step - loss: 0.7718 - a
ccuracy: 0.6598 - val loss: 0.6959 - val accuracy: 0.6862
Epoch 31/80
329/329 [============= ] - 60s 183ms/step - loss: 0.7684 - a
ccuracy: 0.6582 - val loss: 0.6988 - val accuracy: 0.6827
Epoch 32/80
329/329 [============== ] - 60s 183ms/step - loss: 0.7642 - a
ccuracy: 0.6648 - val loss: 0.6940 - val accuracy: 0.6871
Epoch 33/80
329/329 [============= ] - 60s 183ms/step - loss: 0.7668 - a
ccuracy: 0.6617 - val loss: 0.6914 - val accuracy: 0.6871
Epoch 34/80
329/329 [=============== ] - 61s 184ms/step - loss: 0.7603 - a
ccuracy: 0.6619 - val loss: 0.6883 - val accuracy: 0.6898
Epoch 35/80
329/329 [============= ] - 60s 184ms/step - loss: 0.7617 - a
ccuracy: 0.6645 - val loss: 0.6872 - val accuracy: 0.6898
Epoch 36/80
329/329 [============ ] - 60s 183ms/step - loss: 0.7564 - a
ccuracy: 0.6680 - val loss: 0.6864 - val accuracy: 0.6951
Epoch 37/80
```

```
329/329 [=============== ] - 61s 184ms/step - loss: 0.7605 - a
ccuracy: 0.6647 - val loss: 0.6881 - val accuracy: 0.6920
Epoch 38/80
329/329 [============= ] - 60s 182ms/step - loss: 0.7560 - a
ccuracy: 0.6691 - val loss: 0.6842 - val accuracy: 0.6938
Epoch 39/80
329/329 [================== ] - 60s 183ms/step - loss: 0.7549 - a
ccuracy: 0.6688 - val loss: 0.6824 - val accuracy: 0.6920
Epoch 40/80
329/329 [============= ] - 60s 182ms/step - loss: 0.7461 - a
ccuracy: 0.6697 - val loss: 0.6817 - val accuracy: 0.6942
ccuracy: 0.6699 - val loss: 0.6803 - val accuracy: 0.6929
Epoch 42/80
329/329 [============= ] - 60s 183ms/step - loss: 0.7515 - a
ccuracy: 0.6698 - val loss: 0.6808 - val accuracy: 0.6942
Epoch 43/80
ccuracy: 0.6696 - val loss: 0.6800 - val accuracy: 0.6911
Epoch 44/80
329/329 [============ ] - 61s 185ms/step - loss: 0.7524 - a
ccuracy: 0.6659 - val loss: 0.6784 - val accuracy: 0.6942
Epoch 45/80
329/329 [============= ] - 59s 179ms/step - loss: 0.7455 - a
ccuracy: 0.6678 - val loss: 0.6802 - val accuracy: 0.6916
Epoch 46/80
329/329 [============= ] - 59s 180ms/step - loss: 0.7462 - a
ccuracy: 0.6706 - val loss: 0.6801 - val accuracy: 0.6924
Epoch 47/80
329/329 [============ ] - 60s 183ms/step - loss: 0.7423 - a
ccuracy: 0.6724 - val loss: 0.6778 - val accuracy: 0.6933
Epoch 48/80
329/329 [============ ] - 60s 183ms/step - loss: 0.7449 - a
ccuracy: 0.6716 - val loss: 0.6793 - val accuracy: 0.6933
Epoch 49/80
329/329 [============= ] - 60s 183ms/step - loss: 0.7469 - a
ccuracy: 0.6712 - val loss: 0.6767 - val accuracy: 0.6951
Epoch 50/80
329/329 [============ ] - 60s 184ms/step - loss: 0.7450 - a
ccuracy: 0.6673 - val loss: 0.6740 - val accuracy: 0.6991
Epoch 51/80
329/329 [================== ] - 61s 185ms/step - loss: 0.7447 - a
ccuracy: 0.6663 - val_loss: 0.6751 - val_accuracy: 0.6942
Epoch 52/80
329/329 [============= ] - 60s 184ms/step - loss: 0.7434 - a
ccuracy: 0.6723 - val_loss: 0.6757 - val_accuracy: 0.6956
Epoch 53/80
329/329 [============ ] - 61s 184ms/step - loss: 0.7439 - a
ccuracy: 0.6753 - val loss: 0.6743 - val accuracy: 0.6956
Epoch 54/80
329/329 [============ ] - 60s 183ms/step - loss: 0.7377 - a
ccuracy: 0.6696 - val loss: 0.6755 - val accuracy: 0.6951
uracy: 0.6889
Test accuracy: 0.6888889074325562
```

	precision	recall	f1-score	support
Negative Neutral Positive	0.75 0.57 0.77	0.65 0.67 0.74	0.70 0.62 0.76	773 741 736
accuracy macro avg weighted avg	0.70 0.70	0.69 0.69	0.69 0.69 0.69	2250 2250 2250

# Confusion Matrix for Configuration 4

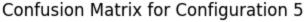


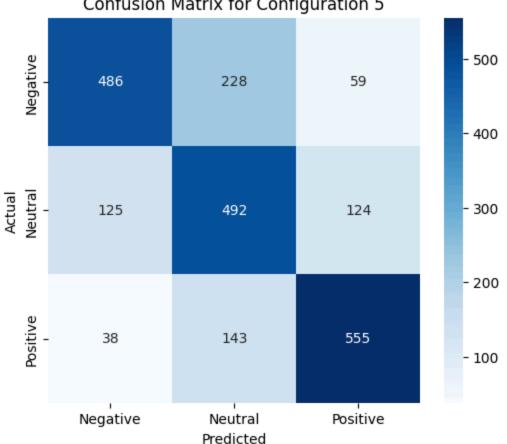
```
New best model found!
Training configuration 5/6: {'optimizer': 'adam', 'learning rate': 1e-05, 'b
atch size': 16, 'dropout rate': 0.1, 'activation function': 'tanh', 'use bat
chnorm': False}
Epoch 1/80
657/657 [===========] - 74s 102ms/step - loss: 1.1318 - a
ccuracy: 0.3467 - val loss: 1.0639 - val accuracy: 0.4324
Epoch 2/80
657/657 [============= ] - 65s 99ms/step - loss: 1.0935 - ac
curacy: 0.3970 - val loss: 1.0220 - val accuracy: 0.5209
Epoch 3/80
curacy: 0.4506 - val loss: 0.9673 - val accuracy: 0.6080
Epoch 4/80
curacy: 0.4947 - val loss: 0.9150 - val accuracy: 0.6133
657/657 [============] - 65s 99ms/step - loss: 0.9579 - ac
curacy: 0.5386 - val loss: 0.8674 - val accuracy: 0.6316
Epoch 6/80
657/657 [============] - 65s 98ms/step - loss: 0.9182 - ac
curacy: 0.5549 - val loss: 0.8318 - val accuracy: 0.6320
Epoch 7/80
curacy: 0.5826 - val loss: 0.8031 - val accuracy: 0.6404
Epoch 8/80
curacy: 0.5877 - val loss: 0.7815 - val accuracy: 0.6516
Epoch 9/80
curacy: 0.5963 - val loss: 0.7684 - val accuracy: 0.6520
Epoch 10/80
657/657 [============== ] - 65s 98ms/step - loss: 0.8403 - ac
curacy: 0.6081 - val loss: 0.7611 - val accuracy: 0.6551
Epoch 11/80
curacy: 0.6164 - val loss: 0.7481 - val accuracy: 0.6622
Epoch 12/80
curacy: 0.6254 - val loss: 0.7434 - val accuracy: 0.6698
Epoch 13/80
curacy: 0.6304 - val loss: 0.7380 - val accuracy: 0.6676
Epoch 14/80
curacy: 0.6256 - val loss: 0.7280 - val accuracy: 0.6702
Epoch 15/80
curacy: 0.6350 - val loss: 0.7249 - val accuracy: 0.6733
Epoch 16/80
curacy: 0.6378 - val loss: 0.7227 - val accuracy: 0.6742
Epoch 17/80
657/657 [===========] - 65s 99ms/step - loss: 0.7997 - ac
curacy: 0.6364 - val loss: 0.7159 - val accuracy: 0.6711
Epoch 18/80
```

```
curacy: 0.6381 - val loss: 0.7161 - val accuracy: 0.6782
Epoch 19/80
curacy: 0.6425 - val loss: 0.7151 - val accuracy: 0.6791
Epoch 20/80
curacy: 0.6433 - val loss: 0.7159 - val accuracy: 0.6791
Epoch 21/80
curacy: 0.6511 - val loss: 0.7079 - val accuracy: 0.6804
curacy: 0.6488 - val loss: 0.7048 - val accuracy: 0.6827
Epoch 23/80
curacy: 0.6495 - val loss: 0.7050 - val accuracy: 0.6831
Epoch 24/80
curacy: 0.6457 - val loss: 0.7057 - val accuracy: 0.6813
Epoch 25/80
curacy: 0.6576 - val loss: 0.6994 - val accuracy: 0.6884
Epoch 26/80
curacy: 0.6595 - val loss: 0.6986 - val accuracy: 0.6876
Epoch 27/80
curacy: 0.6612 - val loss: 0.6952 - val_accuracy: 0.6920
Epoch 28/80
657/657 [============= ] - 64s 98ms/step - loss: 0.7725 - ac
curacy: 0.6559 - val loss: 0.6972 - val accuracy: 0.6902
Epoch 29/80
curacy: 0.6599 - val loss: 0.6931 - val accuracy: 0.6898
Epoch 30/80
curacy: 0.6572 - val loss: 0.6993 - val accuracy: 0.6880
Epoch 31/80
657/657 [============] - 64s 98ms/step - loss: 0.7633 - ac
curacy: 0.6553 - val loss: 0.6981 - val accuracy: 0.6893
curacy: 0.6647 - val loss: 0.6918 - val accuracy: 0.6938
Epoch 33/80
curacy: 0.6605 - val loss: 0.6975 - val accuracy: 0.6902
Epoch 34/80
curacy: 0.6647 - val loss: 0.6962 - val accuracy: 0.6880
Epoch 35/80
curacy: 0.6619 - val loss: 0.6916 - val accuracy: 0.6960
Epoch 36/80
curacy: 0.6572 - val loss: 0.6916 - val accuracy: 0.6982
```

```
Epoch 37/80
curacy: 0.6639 - val loss: 0.6864 - val accuracy: 0.6964
curacy: 0.6636 - val loss: 0.6921 - val accuracy: 0.6938
curacy: 0.6630 - val loss: 0.6910 - val accuracy: 0.6973
Epoch 40/80
curacy: 0.6595 - val loss: 0.6853 - val accuracy: 0.6942
Epoch 41/80
curacy: 0.6658 - val loss: 0.6896 - val accuracy: 0.6969
Epoch 42/80
curacy: 0.6648 - val loss: 0.6859 - val accuracy: 0.6982
657/657 [============= ] - 74s 113ms/step - loss: 0.7504 - a
ccuracy: 0.6707 - val loss: 0.6883 - val accuracy: 0.6933
Epoch 44/80
curacy: 0.6627 - val loss: 0.6840 - val accuracy: 0.6947
Epoch 45/80
curacy: 0.6692 - val loss: 0.6856 - val accuracy: 0.7000
Epoch 46/80
curacy: 0.6687 - val_loss: 0.6875 - val_accuracy: 0.6956
Epoch 47/80
curacy: 0.6687 - val loss: 0.6837 - val accuracy: 0.6978
Epoch 48/80
curacy: 0.6710 - val loss: 0.6828 - val accuracy: 0.6978
curacy: 0.6735 - val loss: 0.6828 - val accuracy: 0.6978
Epoch 50/80
curacy: 0.6670 - val loss: 0.6823 - val accuracy: 0.6964
Epoch 51/80
curacy: 0.6614 - val loss: 0.6821 - val accuracy: 0.6982
Epoch 52/80
curacy: 0.6706 - val loss: 0.6783 - val accuracy: 0.6982
Epoch 53/80
657/657 [============] - 64s 98ms/step - loss: 0.7423 - ac
curacy: 0.6695 - val loss: 0.6789 - val accuracy: 0.6973
Epoch 54/80
657/657 [===========] - 65s 99ms/step - loss: 0.7481 - ac
curacy: 0.6723 - val loss: 0.6795 - val accuracy: 0.6982
```

curacy: 0.6697 - val loss: 0.6771 - val accuracy: 0.6982 Epoch 56/80 curacy: 0.6716 - val loss: 0.6792 - val accuracy: 0.7013 Epoch 57/80 curacy: 0.6691 - val loss: 0.6810 - val accuracy: 0.6996 Epoch 58/80 curacy: 0.6745 - val loss: 0.6806 - val accuracy: 0.6991 Epoch 59/80 657/657 [============= ] - 64s 98ms/step - loss: 0.7470 - ac curacy: 0.6707 - val loss: 0.6778 - val accuracy: 0.7013 curacy: 0.6813 Test accuracy: 0.6813333630561829 141/141 [=======] - 12s 74ms/step Classification Report for Configuration 5: precision recall f1-score support 0.68 Negative 0.75 0.63 773 Neutral 0.57 0.66 0.61 741 Positive 0.75 0.75 0.75 736 0.68 accuracy 2250 0.68 macro avq 0.69 0.68 2250 0.69 0.68 2250 weighted avg 0.68



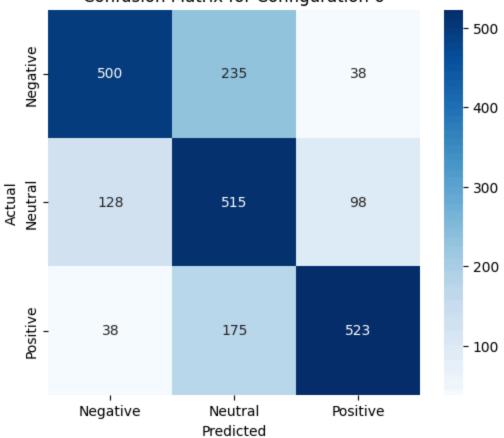


```
Training configuration 6/6: {'optimizer': 'adam', 'learning rate': 3e-05, 'b
atch size': 16, 'dropout rate': 0.2, 'activation function': 'relu', 'use bat
chnorm': False}
Epoch 1/80
657/657 [============= ] - 74s 102ms/step - loss: 1.1243 - a
ccuracy: 0.3430 - val loss: 1.0926 - val accuracy: 0.3338
657/657 [============== ] - 65s 98ms/step - loss: 1.0952 - ac
curacy: 0.3697 - val loss: 1.0708 - val accuracy: 0.4542
Epoch 3/80
curacy: 0.4084 - val loss: 1.0314 - val accuracy: 0.5991
Epoch 4/80
curacy: 0.4702 - val loss: 0.9296 - val accuracy: 0.6027
Epoch 5/80
curacy: 0.5165 - val loss: 0.8441 - val accuracy: 0.6142
curacy: 0.5471 - val loss: 0.8108 - val accuracy: 0.6231
Epoch 7/80
curacy: 0.5716 - val loss: 0.7948 - val accuracy: 0.6347
Epoch 8/80
curacy: 0.5766 - val loss: 0.7734 - val accuracy: 0.6462
Epoch 9/80
curacy: 0.5939 - val loss: 0.7602 - val accuracy: 0.6591
Epoch 10/80
curacy: 0.5951 - val loss: 0.7535 - val accuracy: 0.6627
Epoch 11/80
curacy: 0.6167 - val loss: 0.7391 - val accuracy: 0.6631
curacy: 0.6216 - val loss: 0.7370 - val accuracy: 0.6667
Epoch 13/80
curacy: 0.6295 - val loss: 0.7322 - val accuracy: 0.6707
Epoch 14/80
curacy: 0.6317 - val loss: 0.7274 - val accuracy: 0.6756
Epoch 15/80
curacy: 0.6375 - val loss: 0.7192 - val accuracy: 0.6849
Epoch 16/80
657/657 [============] - 65s 99ms/step - loss: 0.8003 - ac
curacy: 0.6398 - val loss: 0.7176 - val accuracy: 0.6840
Epoch 17/80
curacy: 0.6411 - val loss: 0.7122 - val accuracy: 0.6844
Epoch 18/80
```

```
curacy: 0.6436 - val loss: 0.7108 - val accuracy: 0.6751
Epoch 19/80
curacy: 0.6462 - val loss: 0.7042 - val accuracy: 0.6800
Epoch 20/80
657/657 [============] - 65s 98ms/step - loss: 0.7857 - ac
curacy: 0.6553 - val_loss: 0.7026 - val accuracy: 0.6858
Epoch 21/80
curacy: 0.6583 - val loss: 0.6977 - val accuracy: 0.6849
Epoch 22/80
curacy: 0.6545 - val loss: 0.6994 - val accuracy: 0.6853
Epoch 23/80
curacy: 0.6607 - val loss: 0.6969 - val accuracy: 0.6880
curacy: 0.6663 - val loss: 0.6920 - val accuracy: 0.6907
Epoch 25/80
657/657 [============] - 65s 99ms/step - loss: 0.7592 - ac
curacy: 0.6664 - val loss: 0.6915 - val accuracy: 0.6929
Epoch 26/80
curacy: 0.6643 - val loss: 0.6924 - val accuracy: 0.6911
Epoch 27/80
curacy: 0.6687 - val loss: 0.6864 - val accuracy: 0.6929
Epoch 28/80
curacy: 0.6690 - val loss: 0.6886 - val accuracy: 0.6920
curacy: 0.6694 - val loss: 0.6878 - val accuracy: 0.6929
Epoch 30/80
curacy: 0.6730 - val loss: 0.6766 - val accuracy: 0.6938
Epoch 31/80
curacy: 0.6709 - val loss: 0.6840 - val accuracy: 0.6942
Epoch 32/80
curacy: 0.6677 - val loss: 0.6801 - val accuracy: 0.6920
Epoch 33/80
curacy: 0.6708 - val loss: 0.6787 - val accuracy: 0.6973
Epoch 34/80
curacy: 0.6694 - val loss: 0.6824 - val accuracy: 0.6884
curacy: 0.6836
Test accuracy: 0.683555543422699
Classification Report for Configuration 6:
       precision recall f1-score support
```

Negative	0.75	0.65	0.69	773
Neutral	0.56	0.70	0.62	741
Positive	0.79	0.71	0.75	736
accuracy			0.68	2250
macro avg	0.70	0.68	0.69	2250
weighted avg	0.70	0.68	0.69	2250





### 6.2.2 Non-Batch-Normalized Models: Examples of Mis-classified Points

```
In []: # After the training loop
    print(f"\nBest Model Configuration: {best_config}")
    print(f"Best Model Test Accuracy: {best_accuracy}")

# Use the best model's predictions
    y_pred = best_y_pred
    y_true = best_y_true

# Identify misclassified examples
    misclassified_indices = np.where(y_pred != y_true)[0]

# Decode test texts
    def decode_texts(input_ids):
        return [tokenizer.decode(ids, skip_special_tokens=True) for ids in input
    test_texts = decode_texts(test_input_ids)
```

```
# Extract false positives and false negatives
false positives = []
false negatives = []
# Mapping of label indices to label names
label map = {0: 'Negative', 1: 'Neutral', 2: 'Positive'}
# Loop through misclassified examples to separate false positives and false
for idx in misclassified indices:
   true label = y true[idx]
    predicted label = y pred[idx]
    text = test texts[idx]
    if predicted label == 2 and true label != 2:
        # Model predicted Positive, but true label is Negative or Neutral
        false positives.append((text, label map[true label], label map[predi
    elif predicted label != 2 and true label == 2:
        # Model predicted Negative or Neutral, but true label is Positive
        false negatives append((text, label map[true label], label map[predi
# Display a few examples of false positives
print("\nExamples of False Positives:")
for i in range(min(3, len(false positives))):
    text, true label, predicted label = false positives[i]
    print(f"\nText: {text}")
    print(f"True Label: {true label}")
    print(f"Predicted Label: {predicted label}")
# Display a few examples of false negatives
print("\nExamples of False Negatives:")
for i in range(min(3, len(false negatives))):
    text, true label, predicted label = false negatives[i]
    print(f"\nText: {text}")
    print(f"True Label: {true label}")
    print(f"Predicted Label: {predicted label}")
```

Best Model Configuration: {'optimizer': 'adam', 'learning\_rate': 3e-05, 'bat ch\_size': 32, 'dropout\_rate': 0.2, 'activation\_function': 'relu', 'use\_batch

norm': False}

Best Model Test Accuracy: 0.6888889074325562

Examples of False Positives:

Text: this is quite good ive taken to starting my mornings with a nice hot  $\ensuremath{\mathsf{m}}$ 

ug

True Label: Neutral

Predicted Label: Positive

Text: i love gloria jeans cinnamon nut strudel coffee and decided to try this the flavoring is a little over powering for me if youve tried this flavor before and like it you wont be disappointed with it in the kcups packaging

True Label: Neutral Predicted Label: Positive

Text: we love clif kid z bars i decided to try these based on all of the gre at reviews we bought the strawberry and no one in the house likes them my ki ds are picky so i always try the treats just to see if they are even worth t rying i eat healthy and dont mind most healthy snacks but even i did not car e for these i am going to try another flavor but the strawberry just didnt c ut it here

True Label: Negative Predicted Label: Positive

Examples of False Negatives:

Text: i really only needed the funnelpitcher but it was easy and neat to mak e funnel cakes i had been using a kitchen funnel which made a big mess no matter how careful i was i didnt use the ring because i made the funnel cakes in a small deep fryer the mix was fine but i prefer my own recipe if i could have just bought the funnelpitcher that would have been fine and saved me so me money but im not disappointed

True Label: Positive Predicted Label: Neutral

Text: i was very impressed with the flavor growing up on mrs butterworths i developed a taste for commercial syrup and had a hard time adjusting to the real thing but when i opened up the jug of coombs it was delish i would high ly recommend this product

True Label: Positive Predicted Label: Neutral

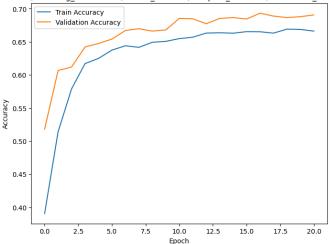
Text: i got a larger one than my puppy probably required and she loves dragg ing it around playing tug of war and just chewing on it unfortunately it doe s shed i dont know if there are rope toys out there that are any sturdier bu t this one does leave little threads behind otherwise great toy

True Label: Positive Predicted Label: Neutral

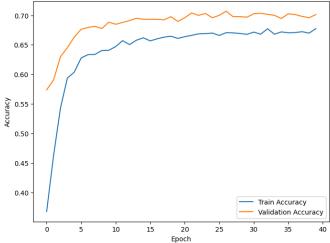
### 6.2.3 Non-Batch-Normalized Models: Visualizations

```
# Plot training and validation accuracy for each configuration
for i, (history, config) in enumerate(histories):
    plt.figure(figsize=(8, 6))
    plt.plot(history.history['accuracy'], label='Train Accuracy')
    plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
    plt.title(f"Configuration {i+1}: {config}")
    plt.xlabel('Epoch')
    plt.ylabel('Accuracy')
    plt.legend()
    plt.show()
```

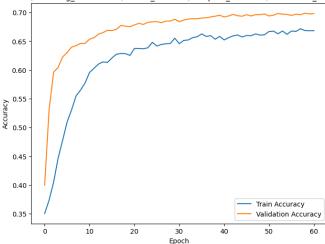
Configuration 1: {'optimizer': 'adamw', 'learning\_rate': 3e-05, 'batch\_size': 16, 'dropout\_rate': 0.1, 'activation\_function': 'tanh', 'use\_batchnorm': False}



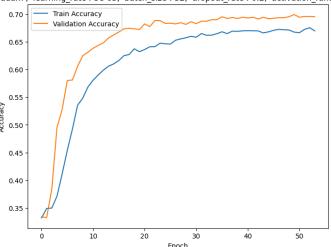
Configuration 2: {'optimizer': 'adamw', 'learning\_rate': 3e-05, 'batch\_size': 32, 'dropout\_rate': 0.1, 'activation\_function': 'tanh', 'use\_batchnorm': False}



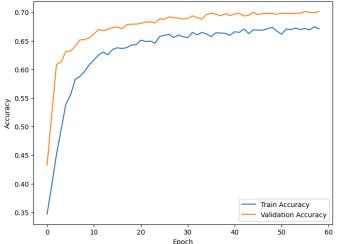
Configuration 3: {'optimizer': 'adamw', 'learning\_rate': 1e-05, 'batch\_size': 32, 'dropout\_rate': 0.1, 'activation\_function': 'tanh', 'use\_batchnorm': False}

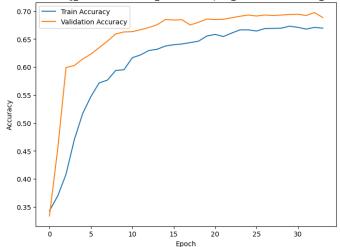


Configuration 4: {'optimizer': 'adam', 'learning\_rate': 3e-05, 'batch\_size': 32, 'dropout\_rate': 0.2, 'activation\_function': 'relu', 'use\_batchnorm': False}

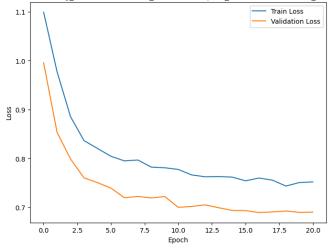


Configuration 5: {'optimizer': 'adam', 'learning\_rate': 1e-05, 'batch\_size': 16, 'dropout\_rate': 0.1, 'activation\_function': 'tanh', 'use\_batchnorm': False}

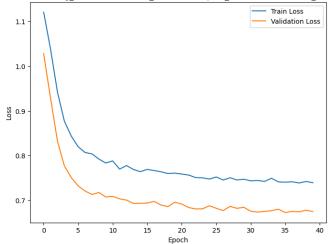




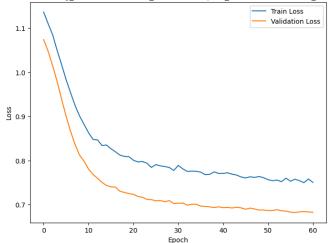
Configuration 1: {'optimizer': 'adamw', 'learning\_rate': 3e-05, 'batch\_size': 16, 'dropout\_rate': 0.1, 'activation\_function': 'tanh', 'use\_batchnorm': False}



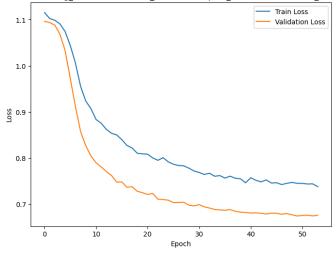
Configuration 2: {'optimizer': 'adamw', 'learning\_rate': 3e-05, 'batch\_size': 32, 'dropout\_rate': 0.1, 'activation\_function': 'tanh', 'use\_batchnorm': False}



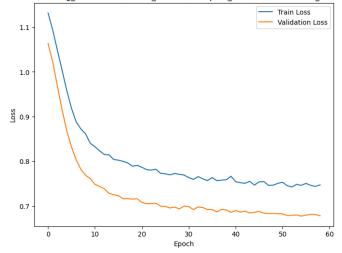
Configuration 3: {'optimizer': 'adamw', 'learning\_rate': 1e-05, 'batch\_size': 32, 'dropout\_rate': 0.1, 'activation\_function': 'tanh', 'use\_batchnorm': False}



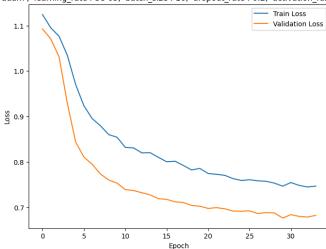
Configuration 4: {'optimizer': 'adam', 'learning\_rate': 3e-05, 'batch\_size': 32, 'dropout\_rate': 0.2, 'activation\_function': 'relu', 'use\_batchnorm': False}



Configuration 5: {'optimizer': 'adam', 'learning\_rate': 1e-05, 'batch\_size': 16, 'dropout\_rate': 0.1, 'activation\_function': 'tanh', 'use\_batchnorm': False}



Configuration 6: {'optimizer': 'adam', 'learning\_rate': 3e-05, 'batch\_size': 16, 'dropout\_rate': 0.2, 'activation\_function': 'relu', 'use\_batchnorm': False}



#### 6.2.4 Non-Batch-Normalized Models: Model Results

```
In []: # Create a DataFrame of the results
    results_df = pd.DataFrame([{
        'optimizer': res['config']['optimizer'],
        'learning_rate': res['config']['learning_rate'],
        'batch_size': res['config']['batch_size'],
        'dropout_rate': res['config']['dropout_rate'],
        'activation_function': res['config']['activation_function'],
        'use_batchnorm': res['config']['use_batchnorm'],
        'accuracy': res['accuracy'],
        'loss': res['loss']
} for res in results])
```

```
optimizer learning rate batch size dropout rate activation function \
                                                  0.1
0
       adamw
                    0.00003
                                     16
1
       adamw
                    0.00003
                                     16
                                                  0.1
                                                                     tanh
2
       adamw
                    0.00003
                                     32
                                                  0.1
                                                                     tanh
                                     32
                                                  0.1
3
       adamw
                    0.00003
                                                                     tanh
4
       adamw
                    0.00001
                                     32
                                                  0.1
                                                                     tanh
5
                                     32
                                                  0.1
       adamw
                    0.00001
                                                                     tanh
6
        adam
                    0.00003
                                     32
                                                  0.2
                                                                     relu
7
                                     32
        adam
                    0.00003
                                                  0.2
                                                                     relu
8
        adam
                    0.00001
                                     16
                                                  0.1
                                                                     tanh
9
        adam
                    0.00001
                                     16
                                                  0.1
                                                                     tanh
                                                  0.2
10
        adam
                    0.00003
                                     16
                                                                     relu
                                     16
                                                  0.2
11
        adam
                    0.00003
                                                                     relu
    use batchnorm accuracy
                                loss
0
            False 0.676000 0.725449
1
            False 0.676000 0.725449
2
            False 0.684444 0.709029
3
            False 0.684444 0.709029
4
            False 0.680444 0.719499
5
            False 0.680444 0.719499
6
            False 0.688889 0.710524
7
            False 0.688889 0.710524
8
            False 0.681333 0.713452
            False 0.681333 0.713452
9
10
            False 0.683556 0.717630
            False 0.683556 0.717630
11
```

## 6.3 Frozen Encoder Layers

### 6.3.1 Frozen Encoder Layers: Training the Model

```
In [ ]: from transformers import TFRobertaForSequenceClassification
        # Load the pre-trained distilroberta model for sequence classification
        model = TFRobertaForSequenceClassification.from pretrained('distilroberta-be
        # Freeze the first 3 encoder layers
        for layer in model.roberta.encoder.layer[:3]:
            layer.trainable = False
        # Define the optimizer
        optimizer = tf.keras.optimizers.Adam(learning rate=3e-5)
        # Define the loss with from logits=True
        loss = tf.keras.losses.SparseCategoricalCrossentropy(from logits=True)
        # Compile the model
        model.compile(
            optimizer=optimizer,
            loss=loss,
            metrics=['accuracy']
        )
        # Set batch size and create datasets
```

```
batch size = 16
train dataset = create tf dataset(train input ids, train attention mask, train
val dataset = create tf dataset(val input ids, val attention mask, val label
test dataset = create tf dataset(test input ids, test attention mask, test l
# Set up callbacks
early stopping = tf.keras.callbacks.EarlyStopping(monitor='val loss', patier
# Fit the model using the validation set
history = model.fit(
   train dataset,
   epochs=10,
   validation data=val dataset,
   callbacks=[early stopping],
   verbose=1
)
# Record the history and config
histories append((history, {'learning rate': 3e-5, 'batch size': 16}))
# Evaluate the final model on the test set
loss, accuracy = model.evaluate(test dataset)
print(f"Test accuracy: {accuracy}")
# Store results
results.append({
    'learning rate': 3e-5,
    'batch size': 16,
    'accuracy': accuracy,
    'loss': loss,
})
# Generate predictions for the test set
y pred probs = model.predict(test dataset)
y pred = np.argmax(y pred probs.logits, axis=1)
# Flatten test labels
y true = test labels
# Compute confusion matrix
cm = confusion matrix(y true, y pred)
# Classification report
report = classification_report(y_true, y_pred, target_names=['Negative', 'Ne
print(f"Classification Report:\n{report}")
# Plot confusion matrix
plt.figure(figsize=(6, 5))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=['Negative',
plt.ylabel('Actual')
plt.xlabel('Predicted')
plt.title(f'Confusion Matrix')
plt.show()
# Store results
results.append({
```

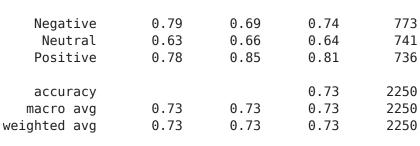
```
'learning rate': 3e-5,
    'batch size': 16,
    'accuracy': accuracy,
    'loss': loss,
    'confusion matrix': cm,
    'classification report': report
})
# If current model is better, save it
# Clear previous best model from memory if exists
best accuracy = accuracy
best model = model
best config = {'learning rate': 3e-5, 'batch size': 16}
best y pred = y pred
best y true = y true
# Also store confusion matrix and classification report
best cm = cm
best report = report
```

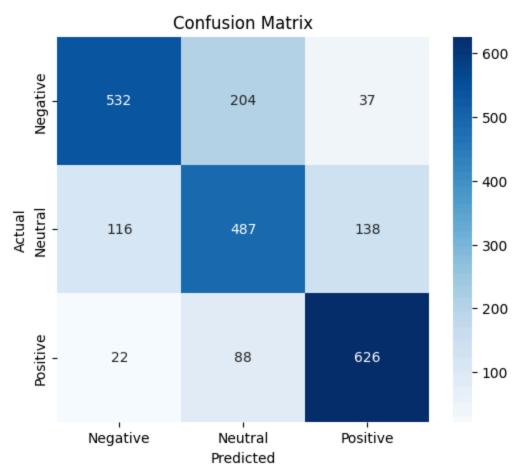
All PyTorch model weights were used when initializing TFRobertaForSequenceCl assification.

Some weights or buffers of the TF 2.0 model TFRobertaForSequenceClassificati on were not initialized from the PyTorch model and are newly initialized: ['classifier.dense.weight', 'classifier.dense.bias', 'classifier.out\_proj.weight', 'classifier.out\_proj.bias']

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

```
Epoch 1/10
accuracy: 0.6799 - val loss: 0.5798 - val accuracy: 0.7462
accuracy: 0.7679 - val loss: 0.6187 - val accuracy: 0.7356
accuracy: 0.8105 - val loss: 0.6324 - val accuracy: 0.7493
Epoch 4/10
accuracy: 0.8486 - val_loss: 0.6656 - val accuracy: 0.7507
curacy: 0.7311
Test accuracy: 0.7311111092567444
141/141 [========== ] - 12s 72ms/step
Classification Report:
       precision recall f1-score
                       support
                    0.74
  Negative
         0.79
               0.69
                          773
```





```
In []: # Save the tokenizer (recommended)
    tokenizer.save_pretrained('model_directory')

# Save the model
    model.save_pretrained('model_directory')
```

#### 6.3.2 Frozen Encoder Layers: Mis-classified Prices

```
In [ ]: # After the training loop
        print(f"\nBest Model Configuration: {best config}")
        print(f"Best Model Test Accuracy: {best accuracy}")
        # Use the best model's predictions
        y pred = best y pred
        y true = best y true
        # Identify misclassified examples
        misclassified_indices = np.where(y_pred != y_true)[0]
        # Decode test texts
        def decode texts(input ids):
            return [tokenizer.decode(ids, skip special tokens=True) for ids in input
        test texts = decode texts(test input ids)
        # Extract false positives and false negatives
        false positives = []
        false negatives = []
        # Mapping of label indices to label names
        label map = {0: 'Negative', 1: 'Neutral', 2: 'Positive'}
        # Loop through misclassified examples to separate false positives and false
        for idx in misclassified indices:
            true label = y_true[idx]
            predicted label = y pred[idx]
            text = test texts[idx]
            if predicted label == 2 and true label != 2:
                # Model predicted Positive, but true label is Negative or Neutral
                false positives append((text, label map[true label], label map[predi
            elif predicted label != 2 and true label == 2:
                # Model predicted Negative or Neutral, but true label is Positive
                false negatives append((text, label map[true label], label map[predi
        # Display a few examples of false positives
        print("\nExamples of False Positives:")
        for i in range(min(3, len(false positives))):
            text, true label, predicted label = false positives[i]
            print(f"\nText: {text}")
            print(f"True Label: {true label}")
            print(f"Predicted Label: {predicted label}")
        # Display a few examples of false negatives
        print("\nExamples of False Negatives:")
```

```
for i in range(min(3, len(false_negatives))):
    text, true_label, predicted_label = false_negatives[i]
    print(f"\nText: {text}")
    print(f"True Label: {true_label}")
    print(f"Predicted Label: {predicted_label}")
```

Best Model Configuration: {'learning\_rate': 3e-05, 'batch\_size': 16} Best Model Test Accuracy: 0.7311111092567444

Examples of False Positives:

Text: this is quite good ive taken to starting my mornings with a nice hot  ${\tt m}$ 

ug

True Label: Neutral

Predicted Label: Positive

Text: i love gloria jeans cinnamon nut strudel coffee and decided to try this the flavoring is a little over powering for me if youve tried this flavor before and like it you wont be disappointed with it in the kcups packaging

True Label: Neutral Predicted Label: Positive

Text: im an espresso nut and have several machines that i call my friends on e of these is a jura s avantgarde and after trying dozens and dozens of various beans to find the perfect mate intelligentsia black cat classic prevaile d as the very clear winner there may be better beans out there but not for t his machine ive served thousands of shots of this espresso out of the jura a nd everyone is impressed by it actually not just impressed blown away its not uncommon to have neighbors coming in the back door in the morning to grab a shot this bean produces a perfect espresso great crema thick legs smooth w ith no bitterness and a slight nut

True Label: Neutral Predicted Label: Positive

Examples of False Negatives:

Text: i really only needed the funnelpitcher but it was easy and neat to mak e funnel cakes i had been using a kitchen funnel which made a big mess no matter how careful i was i didnt use the ring because i made the funnel cakes in a small deep fryer the mix was fine but i prefer my own recipe if i could have just bought the funnelpitcher that would have been fine and saved me so me money but im not disappointed

True Label: Positive Predicted Label: Neutral

Text: i got a larger one than my puppy probably required and she loves dragg ing it around playing tug of war and just chewing on it unfortunately it does shed i dont know if there are rope toys out there that are any sturdier but this one does leave little threads behind otherwise great toy

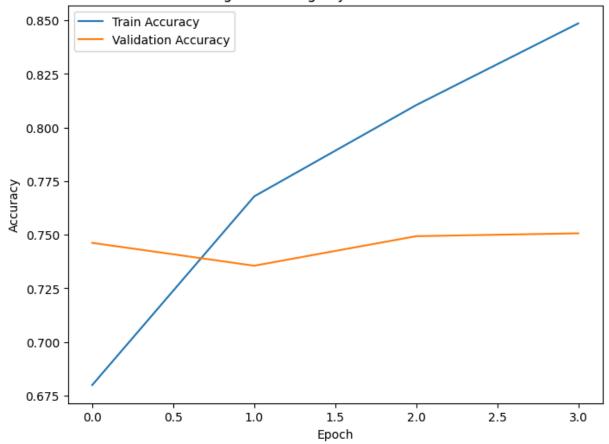
True Label: Positive Predicted Label: Neutral

Text: our dogs love dingo bones but these peanut butter flavor bones are an obsession of our smallest rat terrier hes only pounds so bones tend to last several sittings for him give him one of these and hes occupied for as long as it takes to finish hell still play with other bones but he loses his mind when he sees me go for this bag i just wish they were priced as well as the other dingo products and less difficult to find

True Label: Positive Predicted Label: Neutral

```
In []: # Plot training and validation accuracy for each configuration
for i, (history, config) in enumerate(histories):
    if i == 6:
        plt.figure(figsize=(8, 6))
        plt.plot(history.history['accuracy'], label='Train Accuracy')
        plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
        plt.title("Fine-Tuning 3 Encoding Layers of distilroberta-base")
        plt.xlabel('Epoch')
        plt.ylabel('Accuracy')
        plt.legend()
        plt.show()
    else:
        None
```





As seen above, this model has the highest test accuracy and as such is our best performing model. This will be the model we use for the inference pipeline.

# 7. Inference Pipeline

Below is the inference pipeline for our model. It is a function that takes text and returns the predicted sentiment.

```
In [ ]: from transformers import RobertaTokenizer, TFRobertaForSequenceClassification
        # Load the tokenizer
        tokenizer = RobertaTokenizer.from pretrained('model directory')
        # Load the model
        best model = TFRobertaForSequenceClassification.from pretrained('model direct
In [ ]: import re
        def clean text(text):
            # Remove HTML tags
            text = re.sub(r'<.*?>', '', text)
            # Remove non-ASCII characters
            text = text.encode('ascii', 'ignore').decode('utf-8')
            # Remove URLs
            text = re.sub(r'http\S+', '', text)
            # Remove special characters and numbers
            text = re.sub(r'[^A-Za-z\s]', '', text)
            # Convert to lowercase
            text = text.lower()
            # Remove extra spaces
            text = re.sub(r'\s+', ' ', text).strip()
            return text
In [ ]: import tensorflow as tf
        def predict sentiment(text):
            # Clean the input text
            cleaned text = clean text(text)
            # Tokenize the text
            inputs = tokenizer(
                cleaned text,
                padding='max length',
                truncation=True,
                max length=128,
                return tensors='tf'
            )
            # Get model predictions
            outputs = best model(inputs)
            logits = outputs.logits
            # Get the predicted class
            predicted class = tf.argmax(logits, axis=1).numpy()[0]
            # Map the predicted class to sentiment label
            sentiment_labels = {0: 'Negative', 1: 'Neutral', 2: 'Positive'}
            predicted sentiment = sentiment labels[predicted class]
            return predicted sentiment
        # Example usage
        sample text = "This product is fantastic! I absolutely love it."
```

```
predicted_sentiment = predict_sentiment(sample_text)
print(f"\nText: {sample_text}")
print(f"Predicted Sentiment: {predicted_sentiment}")
```

All model checkpoint layers were used when initializing TFRobertaForSequence Classification.

All the layers of TFRobertaForSequenceClassification were initialized from the model checkpoint at model\_directory.

If your task is similar to the task the model of the checkpoint was trained on, you can already use TFRobertaForSequenceClassification for predictions w ithout further training.

Text: This product is fantastic! I absolutely love it.

Predicted Sentiment: Positive

## 8. Reflections

Andre: Creating these models was very eye opening for me. In particular, I learned a lot about how long it takes in order to properly train models. Along with this, I learned about the importance of the quantity of training data for fine tuning models. Finally, I learned about the importance of saving models having lost 2 models that took upwards of an hour to train each. Having done this project, I am now far more comfortable with google colab and how to train models using GPUs.

Adnan: After preprocessing the data and ensuring that each class had an equal number of samples, I used the pre-trained model "cardiffnlp/twitter-roberta-basesentiment-latest" and applied the model's tokenizer for tokenization. I used the full set of 127,920 samples, with 42,640 samples for each class: Negative, Neutral, and Positive. I utilized the Trainer and set the following parameters in the TrainingArguments: output directory as "./results", evaluation strategy as "steps", with evaluations and model saves every 450 steps, a reduced batch size of 16 for both training and evaluation, 4 training epochs, a learning rate of 2e-5, weight decay of 0.01, seed set to 42, logging every 50 steps, and mixed precision enabled (fp16). Additionally, gradient accumulation steps were set to 4, simulating a larger batch size, and tensorboard was used for logging. The results I obtained were: for Class 0 (Negative), Precision: 0.85, Recall: 0.84, F1-Score: 0.85; for Class 1 (Neutral), Precision: 0.78, Recall: 0.80, F1-Score: 0.79; and for Class 2 (Positive), Precision: 0.92, Recall: 0.91, F1-Score: 0.92. The overall Precision for each class was [0.8547, 0.7764, 0.9239], Recall was [0.8409, 0.8036, 0.9063], and F1-Score was [0.8478, 0.7898, 0.9151]. I wanted to experiment with changes in batch size, learning rate, and the number of epochs, as well as trying a manual optimizer with grid search. However, the training took 6 hours and 30 minutes, and I lacked the time and resources to implement these changes.

Majed: Training these models taught me the importance of resource management, early in the model training, I ended up using all 100 units over one night to train a model on the entire dataset. Additionally, I learned that brute force is never the answer, while training on the entire dataset might be enticing, we should train on a portion of the dataset until we reach a desirable model with proven potential to perform well using the entire dataset. Lastly, while this work isn't in the final report, it is on the github, I learned how to create a pipeline which can intake raw data, and convert it through multiple steps to a desirable and usable output.

Yonotan: In my recent work, I gained valuable insights into the use of LSTM models for sentiment analysis, particularly in handling sequential data like product reviews. Implementing GloVe embeddings enhanced my understanding of pre-trained word embeddings and their role in improving model performance while reducing training time. I also tackled the challenge of data imbalance, learning the importance of addressing skewed label distributions to avoid biased models. Through model evaluation, I observed the differences in accuracy between validation and testing datasets, emphasizing the need for robust evaluation to ensure generalizability. Moving forward, I aim to focus on tuning hyperparameters such as learning rates, dropout rates, and epochs to further optimize performance. Additionally, a deeper analysis of misclassified instances will provide valuable insights into areas where the model struggles, allowing for targeted improvements.

Younes: We applied the pre-trained model from hugging face roBERTa, version "roberta-base". We reduced the data as we have limited GPU resources.

Description of hyperparameters: Batch size = 16 Max length Tokenization = 128

Learning\_rate=1e-6 Optimizer= Adam Loss function = Sparse Categorical Crossentropy Number of epochs = 10

We performed 3 trainings: -First training with 3000 samples resulted in underfitting, (Test loss: 58.82% - Test Accuracy: 76%) -Second training with 30000 samples resulted in overfitting, (Test loss: 56.45% - Test Accuracy: 76.32%) -Third training with 30000 samples resulted in overfitting too, but a slight improvement (Test loss: 55.26% - Test Accuracy: 77.23%)

To improve, we can tune the hyperparameters by increasing the size of batches or adding more epochs. We can also try transfer learning by freezing the first layers, and see how the model can behave by training the last layers. We can also try other learning rates.